

# Hiding Access Patterns

## Obliviousness and Differential Privacy

Giuseppe Persiano

Università di Salerno

February 12, 2019

Describing joint work with:  
Sarvar Patel, Mariana Raykova and Kevin Yeo (Google LLC)

- 1 Privacy in Cloud Storage
- 2 Oblivious Algorithms
- 3 An inefficient ORAM
- 4 An insecure ORAM
- 5 A first secure ORAM
- 6 Shuffling without Sorting
- 7 A second construction
- 8 A Recursive Construction
- 9 Lower Bounds
- 10 Differential Privacy
- 11 Where are we?

# Cloud Storage (simplified)

## The perfect marriage of two parties

- The Data Owner  $\mathcal{O}$ :  
owns large amount of data and not enough local storage
- The Storage Manager  $\mathcal{M}$ :  
owns large amount of storage and not enough data

# Cloud Storage (simplified)

## The perfect marriage of two parties

- The Data Owner  $\mathcal{O}$ :  
owns large amount of data and not enough local storage
- The Storage Manager  $\mathcal{M}$ :  
owns large amount of storage and not enough data

If  $\mathcal{O}$  and  $\mathcal{M}$  trust each other

# Cloud Storage (simplified)

## The perfect marriage of two parties

- The Data Owner  $\mathcal{O}$ :  
owns large amount of data and not enough local storage
- The Storage Manager  $\mathcal{M}$ :  
owns large amount of storage and not enough data

If  $\mathcal{O}$  and  $\mathcal{M}$  trust each other

no problem. we can go home now.

# Cloud Storage (simplified)

## The perfect marriage of two parties

- The Data Owner  $\mathcal{O}$ :  
owns large amount of data and not enough local storage
- The Storage Manager  $\mathcal{M}$ :  
owns large amount of storage and not enough data

If  $\mathcal{O}$  and  $\mathcal{M}$  trust each other

no problem. we can go home now.

Lack of trust is much more interesting.

# Enter Encryption

$\mathcal{O}$  does not trust  $\mathcal{M}$  because  $\mathcal{O}$ 's data contain personal data.

# Enter Encryption

$\mathcal{O}$  should not trust  $\mathcal{M}$  because  $\mathcal{O}$ 's data contain personal data.



# Enter Encryption

$\mathcal{O}$  should not trust  $\mathcal{M}$  because  $\mathcal{O}$ 's data contain personal data.

## Use Encryption

- **Private Key:** if  $\mathcal{O}$  is the source of data
- **Public Key:** if data come from various sources

# Enter Encryption

$\mathcal{O}$  should not trust  $\mathcal{M}$  because  $\mathcal{O}$ 's data contain personal data.

## Use Encryption

- **Private Key:** if  $\mathcal{O}$  is the source of data
- **Public Key:** if data come from various sources

Data is

- encrypted before being uploaded to  $\mathcal{M}$

# Enter Encryption

$\mathcal{O}$  should not trust  $\mathcal{M}$  because  $\mathcal{O}$ 's data contain personal data.

## Use Encryption

- **Private Key:** if  $\mathcal{O}$  is the source of data
- **Public Key:** if data come from various sources

## Data is

- encrypted before being uploaded to  $\mathcal{M}$
- decrypted when downloaded from  $\mathcal{M}$

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

## Example

4 customers must be sorted according to revenue.

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

## Example

4 customers must be sorted according to revenue.

- 1 download 1 and 3. decrypt, swap if out of order, re-encrypt, upload.

C;100

D;150

A;200

B;300



# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

## Example

4 customers must be sorted according to revenue.

- 1 download 1 and 3. decrypt, swap if out of order, re-encrypt, upload.
- 2 download 2 and 4. decrypt, swap if out of order, re-encrypt, upload.

C;100

D;150

A;200

B;300

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

## Example

4 customers must be sorted according to revenue.

- 1 download 1 and 3. decrypt, swap if out of order, re-encrypt, upload.
- 2 download 2 and 4. decrypt, swap if out of order, re-encrypt, upload.
- 3 download 1 and 2. decrypt, swap if out of order, re-encrypt, upload.

C;100

D;150

A;200

B;300

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

## Example

4 customers must be sorted according to revenue.

- 1 download 1 and 3. decrypt, swap if out of order, re-encrypt, upload.
- 2 download 2 and 4. decrypt, swap if out of order, re-encrypt, upload.
- 3 download 1 and 2. decrypt, swap if out of order, re-encrypt, upload.
- 4 download 3 and 4. decrypt, swap if out of order, re-encrypt, upload.

C;100

D;150

A;200

B;300

# Are we done?

What if  $\mathcal{O}$  wants to run an algorithm on the encrypted data?

Running an algorithm might reveal information on the data.

Suppose  $\mathcal{O}$  wants to sort the data.

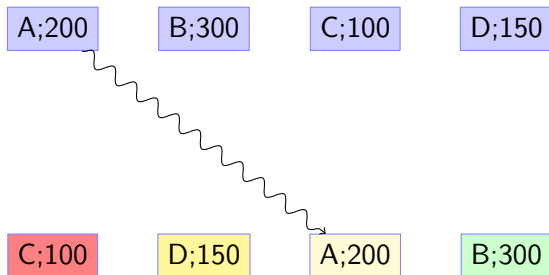
## Example

4 customers must be sorted according to revenue.

- ① download 1 and 3. decrypt, swap if out of order, re-encrypt, upload.
- ② download 2 and 4. decrypt, swap if out of order, re-encrypt, upload.
- ③ download 1 and 2. decrypt, swap if out of order, re-encrypt, upload.
- ④ download 3 and 4. decrypt, swap if out of order, re-encrypt, upload.
- ⑤ download 2 and 3. decrypt, swap if out of order, re-encrypt, upload.

# Security

Can  $\mathcal{M}$  link the first record in the starting configuration to its position in the last configuration?



# Two Concepts

## Indistinguishability of *Swap or Not*

- Download, Decrypt, **Swap or Not**, Re-encrypt, Upload

# Two Concepts

## Indistinguishability of *Swap or Not*

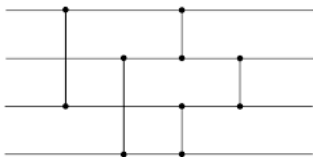
- Download, Decrypt, **Swap or Not**, Re-encrypt, Upload

**Chosen-Ciphertext Security:** Standard notion of security for encryption guarantee that  $\mathcal{M}$  is unable to deduce if a swap has happened.

# Enter Obliviousness

## Definition (Weak Obliviousness)

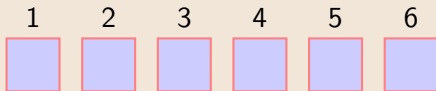
An algorithm is *weakly oblivious* if the *access pattern* to data is the same for all possible inputs of the same length.



Thanks to Wikipedia for the image



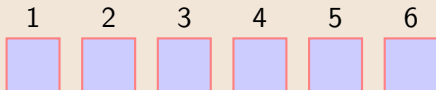
# The adversarial setting



$\mathcal{M}$   
 $\mathcal{O}$

A horizontal dashed green line extends from the vertical position of the  $\mathcal{O}$  symbol across the slide.

# The adversarial setting

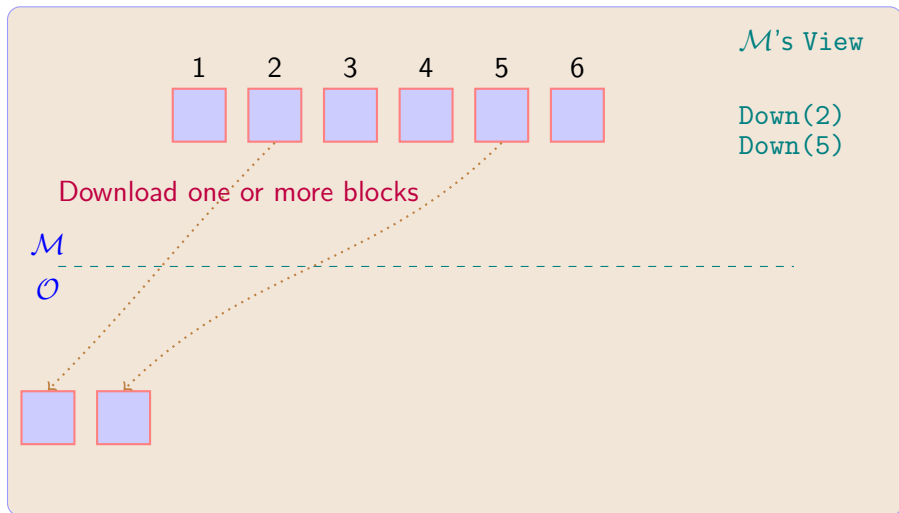


Download one or more blocks

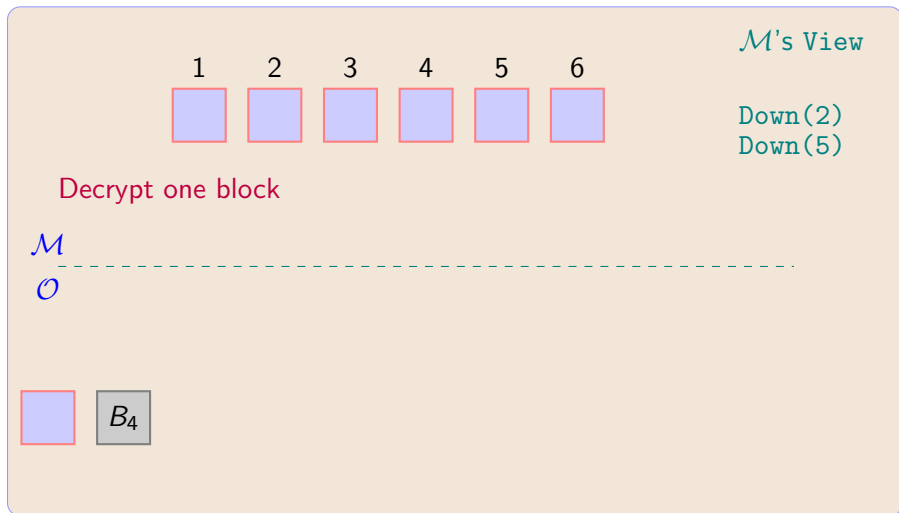
$\mathcal{M}$

$\mathcal{O}$

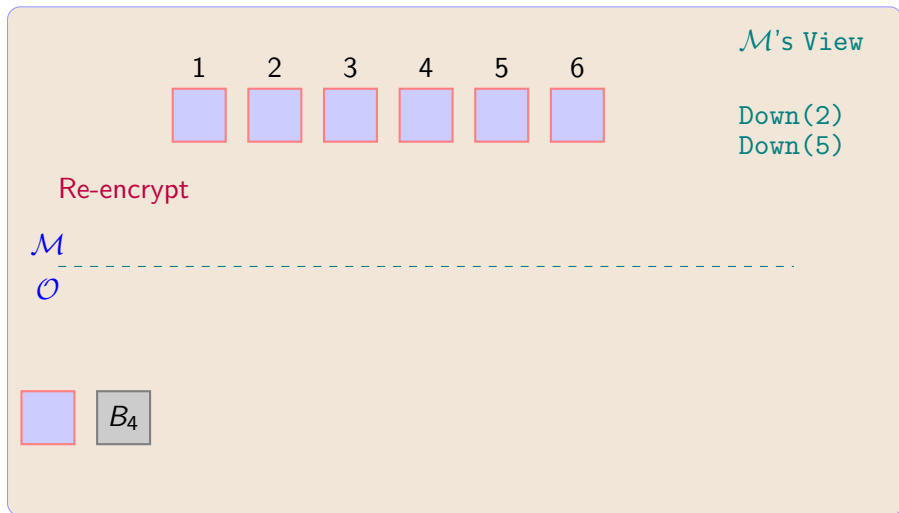
# The adversarial setting



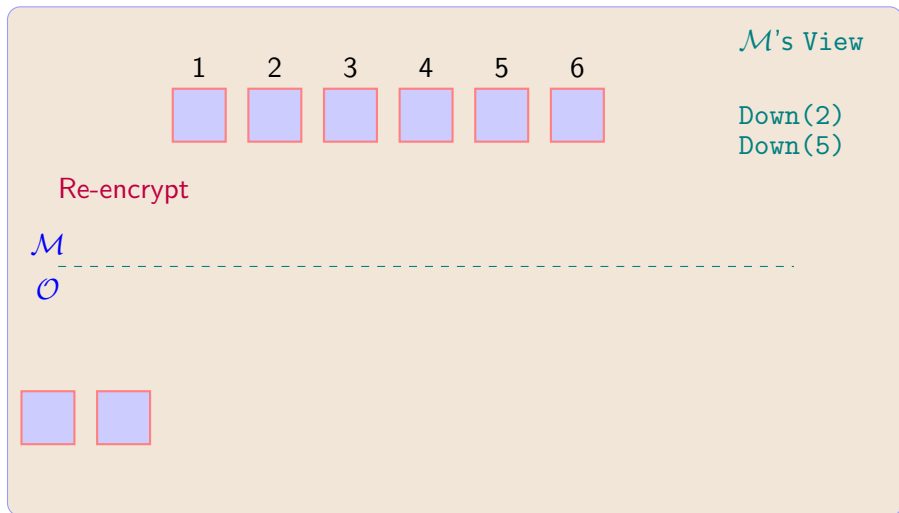
# The adversarial setting



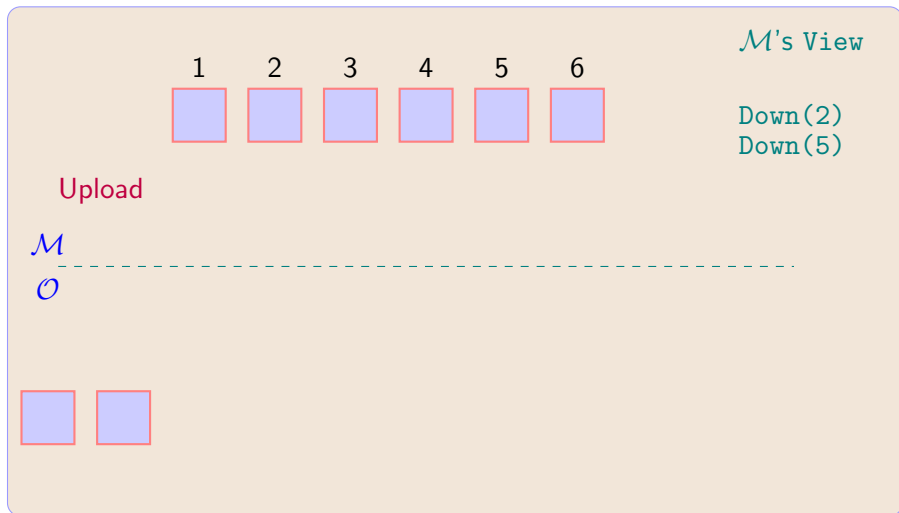
# The adversarial setting



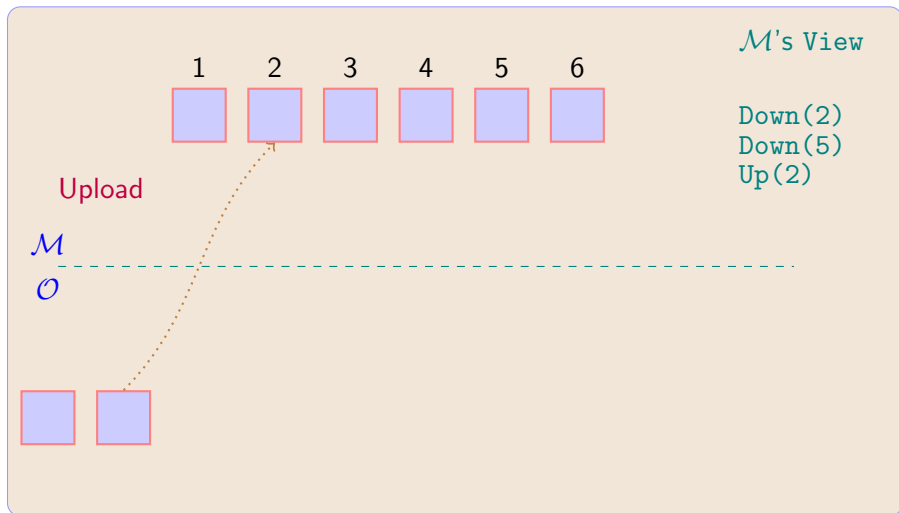
# The adversarial setting



# The adversarial setting



# The adversarial setting





# A new industry

## Job Opportunities for Algorithmists

- Re-design all algorithms to be oblivious!
- Remove all *ifs*, and *whiles*
- **Insertion Sort is not oblivious:**
  - ▶ when the last element of the array is inserted,  $\mathcal{M}$  sees where it lands

# A new industry

## Job Opportunities for Algorithmists

- Re-design all algorithms to be oblivious!
- Remove all *ifs*, and *whiles*
- Insertion Sort is not oblivious:
  - ▶ when the last element of the array is inserted,  $\mathcal{M}$  sees where it lands

Abbiamo abolito la povertà

# A new industry

## Job Opportunities for Algorithmists

- Re-design all algorithms to be oblivious!
- Remove all *ifs*, and *whiles*
- Insertion Sort is not oblivious:
  - ▶ when the last element of the array is inserted,  $\mathcal{M}$  sees where it lands

Abbiamo abolito la povertà ... per gli algoritmist

# A new industry

## Job Opportunities for Algorithmists

- Re-design all algorithms to be oblivious!
- Remove all *ifs*, and *whiles*
- Insertion Sort is not oblivious:
  - ▶ when the last element of the array is inserted,  $\mathcal{M}$  sees where it lands

What? Just move on to the next slide and stop talking politics

# Hiding the Algorithm

## A new threat

- which algorithm is being run **should** also be private information

A;200

B;300

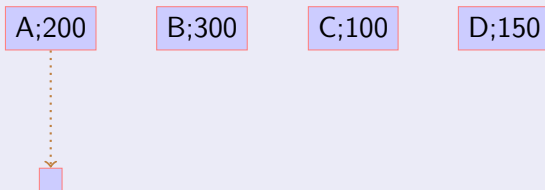
C;100

D;150

# Hiding the Algorithm

## A new threat

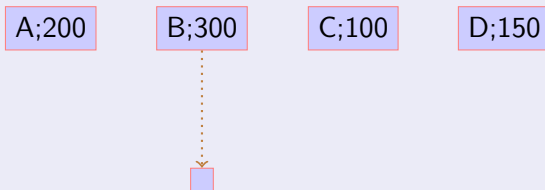
- which algorithm is being run **should** also be private information



# Hiding the Algorithm

## A new threat

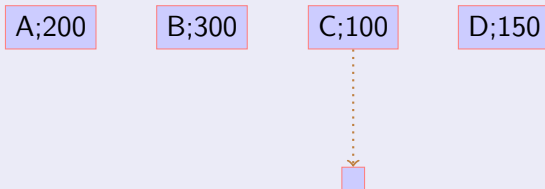
- which algorithm is being run **should** also be private information



# Hiding the Algorithm

## A new threat

- which algorithm is being run **should** also be private information

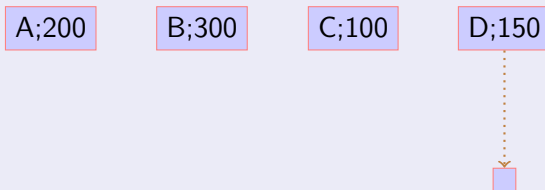




# Hiding the Algorithm

## A new threat

- which algorithm is being run **should** also be private information



## ORAM [Goldreich-Ostrovsky]

- $\mathcal{M}$  stores  $n$  blocks of memory.
  - Every time  $\mathcal{O}$  wants a block, he asks  $\mathcal{M}$  one or more blocks.
  - Security notion:
    - ▶ For any two block sequences  $\mathbb{B} = B_1, \dots, B_n$  and  $\mathbb{C} = C_1, \dots, C_n$
    - ▶ For any two access sequences  $I = (i_1, \dots, i_l)$  and  $J = (j_1, \dots, j_l)$ 
      - ★ performing accesses  $i_1, \dots, i_l$  on  $\mathbb{B} = B_1, \dots, B_n$ ;
      - ★ performing access  $j_1, \dots, j_l$  on  $\mathbb{C} = C_1, \dots, C_n$
- generate the same distribution of accesses to the data stored by  $\mathcal{M}$

## ORAM [Goldreich-Ostrovsky]

- $\mathcal{M}$  stores  $n$  blocks of memory.
  - Every time  $\mathcal{O}$  wants a block, he asks  $\mathcal{M}$  one or more blocks.
  - Security notion:
    - ▶ For any two block sequences  $\mathbb{B} = B_1, \dots, B_n$  and  $\mathbb{C} = C_1, \dots, C_n$
    - ▶ For any two access sequences  $I = (i_1, \dots, i_l)$  and  $J = (j_1, \dots, j_l)$ 
      - ★ performing accesses  $i_1, \dots, i_l$  on  $\mathbb{B} = B_1, \dots, B_n$ ;
      - ★ performing access  $j_1, \dots, j_l$  on  $\mathbb{C} = C_1, \dots, C_n$
- generate the same distribution of accesses to the data stored by  $\mathcal{M}$

For every predicate  $A$

$$\begin{aligned} & \text{Prob}[\text{view} \leftarrow \text{View}(I, \mathbb{B}) : A(\text{view}) = 1] \\ & \leq e^0 \cdot \text{Prob}[\text{view} \leftarrow \text{View}(J, \mathbb{C}) : A(\text{view}) = 1] + \text{negl}(n) \end{aligned}$$

# ORAM makes all Algorithms Oblivious

## Composing ORAM and Non-Oblivious Algorithms

- $\mathcal{O}$  runs the algorithm
- when a block of memory is requested,  $\mathcal{O}$  retrieves it from  $\mathcal{M}$  using ORAM.

# ORAM makes all Algorithms Oblivious

## Composing ORAM and Non-Oblivious Algorithms

- $\mathcal{O}$  runs the algorithm
- when a block of memory is requested,  $\mathcal{O}$  retrieves it from  $\mathcal{M}$  using ORAM.

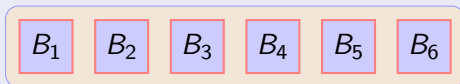
Is ORAM possible at all?

# Yes! This is possible!

## A Trivial ORAM

[▶ Jump ahead](#)

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



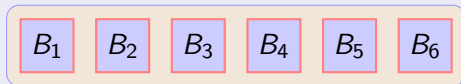
- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

# Yes! This is possible!

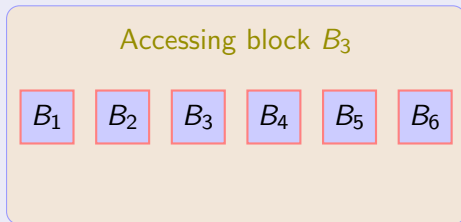
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

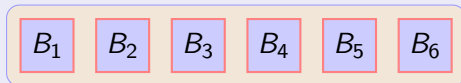


# Yes! This is possible!

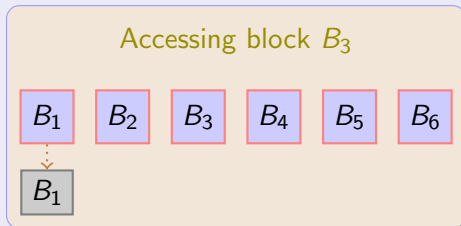
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.



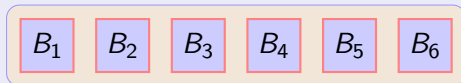


# Yes! This is possible!

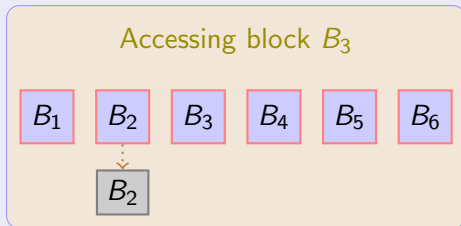
## A Trivial ORAM

[▶ Jump ahead](#)

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

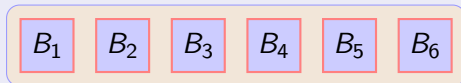


# Yes! This is possible!

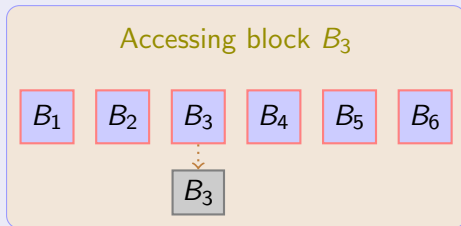
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

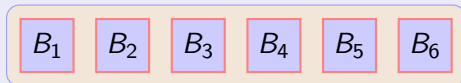


# Yes! This is possible!

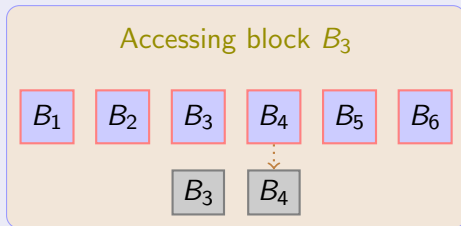
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

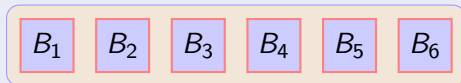


# Yes! This is possible!

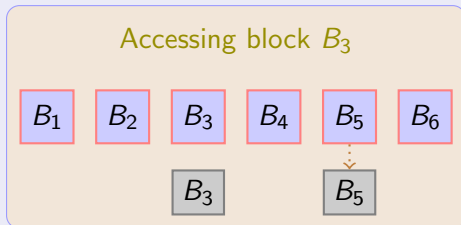
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

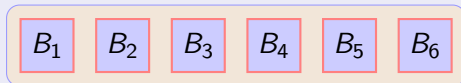


# Yes! This is possible!

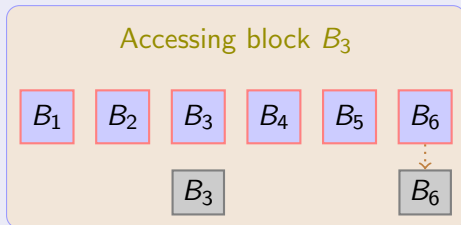
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.

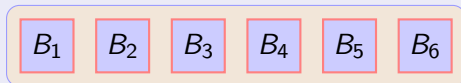


# Yes! This is possible!

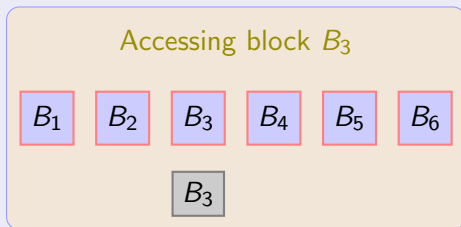
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.



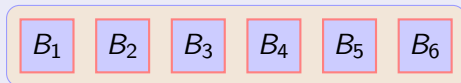
**Access pattern independent from the block accessed  
but...**

# Yes! This is possible!

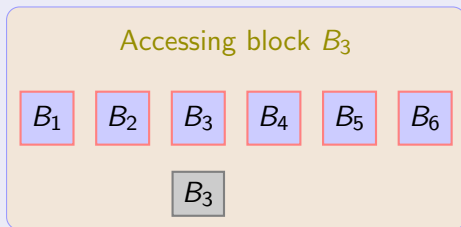
## A Trivial ORAM

► Jump ahead

- All blocks are uploaded to  $\mathcal{M}$  in encrypted form.



- Every time  $\mathcal{O}$  needs to access block  $B_i$ , all the blocks are downloaded and all except for  $B_i$  are discarded.



**Access pattern independent from the block accessed  
but...**

**linear slowdown!!**

First try

**Can this be made efficient?**



# First try

Can this be made efficient?

## First try: Initialization

- permute blocks according to permutation  $\pi$ 
  - ▶ an encryption of  $B_i$  is uploaded in position  $\pi(i)$ ;



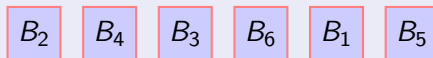
- $\mathcal{O}$  keeps  $\pi$  private;

# First try

Can this be made efficient?

## First try: Initialization

- permute blocks according to permutation  $\pi$ 
  - ▶ an encryption of  $B_i$  is uploaded in position  $\pi(i)$ ;



- $\mathcal{O}$  keeps  $\pi$  private;

# First try

## Can this be made efficient?

### First try: Reading block $i$

- ask  $\mathcal{M}$  for block in position  $\pi(i)$ ;
- decrypt to obtain  $B_i$ ;
- re-encrypt and upload in position  $\pi(i)$ ;

Accessing block  $B_3$

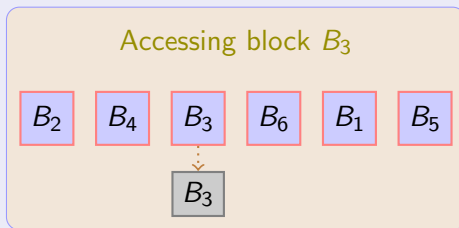


# First try

## Can this be made efficient?

### First try: Reading block $i$

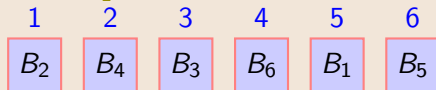
- ask  $\mathcal{M}$  for block in position  $\pi(i)$ ;
- decrypt to obtain  $B_i$ ;
- re-encrypt and upload in position  $\pi(i)$ ;



# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_1$

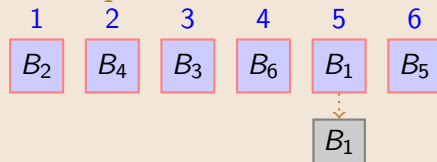


Access pattern seen by  $\mathcal{M}$  :

# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_1$

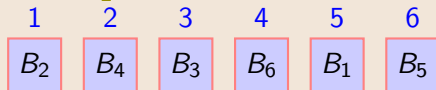


Access pattern seen by  $\mathcal{M}$  : 5

# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_2$

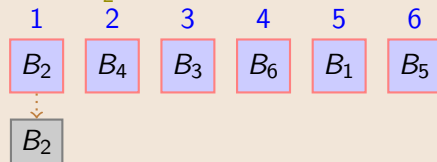


Access pattern seen by  $\mathcal{M}$  : 5

# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_2$



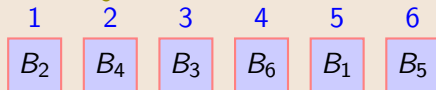
Access pattern seen by  $\mathcal{M}$  : 5, 1



# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_3$

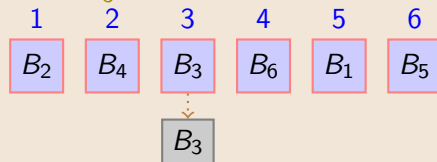


Access pattern seen by  $\mathcal{M}$  : 5, 1

# First try: Security

Access sequence:  $B_1, B_2, B_3$

Accessing block  $B_3$



Access pattern seen by  $\mathcal{M}$  : 5, 1, 3

# First try: Security

Access sequence:  $B_1, B_2, B_3$

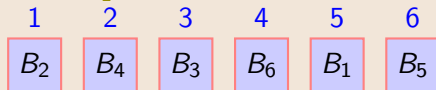
1	2	3	4	5	6
$B_2$	$B_4$	$B_3$	$B_6$	$B_1$	$B_5$

Access pattern seen by  $\mathcal{M}$  :  $x, y, z$

# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_1$

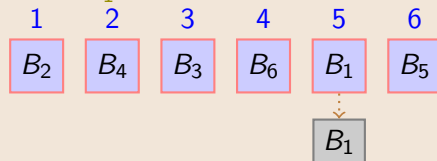


Access pattern seen by  $\mathcal{M}$  :

# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_1$



Access pattern seen by  $\mathcal{M}$  : 5

# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_2$

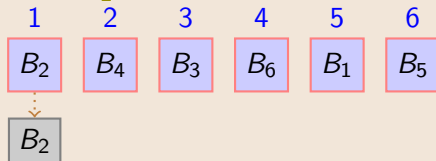


Access pattern seen by  $\mathcal{M}$  : 5

# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_2$



Access pattern seen by  $\mathcal{M}$  : 5, 1

# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_1$



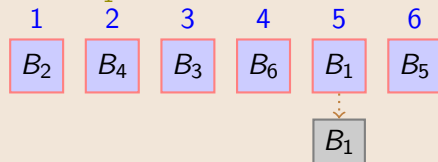
Access pattern seen by  $\mathcal{M}$  : 5, 1



# First try: Security

Access sequence:  $B_1, B_2, B_1$

Accessing block  $B_1$



Access pattern seen by  $\mathcal{M}$  : 5, 1, 5

# First try: Security

Access sequence:  $B_1, B_2, B_1$

1	2	3	4	5	6
$B_2$	$B_4$	$B_3$	$B_6$	$B_1$	$B_5$

Access pattern seen by  $\mathcal{M}$  :  $x, y, x$

# Oblivious RAM

## Obliviousness

For any two *access sequences*  $O_1 = (i_1^1, \dots, i_l^1)$   $O_2 = (i_1^2, \dots, i_l^2)$  of the same length, the distribution of the positions requested by  $\mathcal{O}$  to  $\mathcal{M}$  is the same.

# Oblivious RAM

## Obliviousness

For any two *access sequences*  $O_1 = (i_1^1, \dots, i_l^1)$   $O_2 = (i_1^2, \dots, i_l^2)$  of the same length, the distribution of the positions requested by  $\mathcal{O}$  to  $\mathcal{M}$  is the same.

## Oblivious for Non-repeating sequences

- $k_1 \neq k_2$  implies  $i_{k_1}^1 \neq i_{k_2}^1$  and  $i_{k_1}^2 \neq i_{k_2}^2$ ;
- $\mathcal{M}$  sees requests for  $l$  different randomly chosen blocks both for  $O_1$  and for  $O_2$ .

# Repetition Pattern is leaked

## Repetition Pattern

If the same block is requested twice by  $\mathcal{O}$  then  $\mathcal{M}$  sees the same position accessed twice.

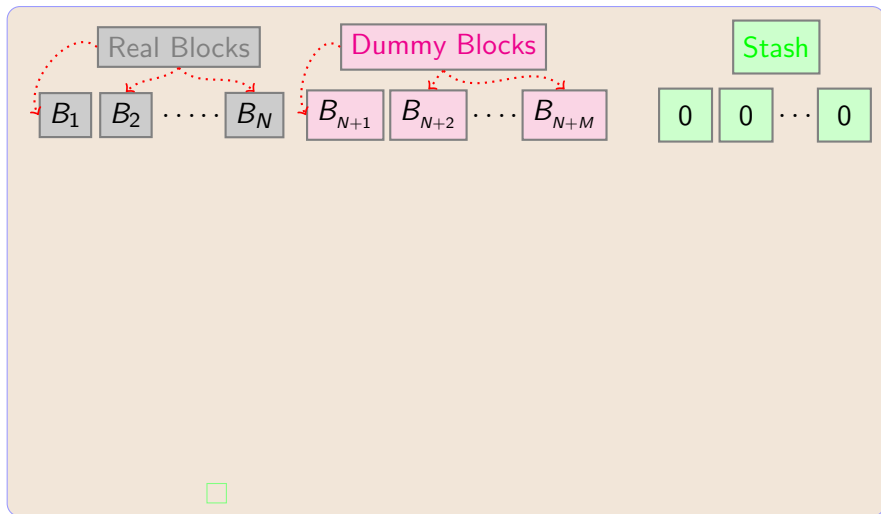
Block	3	4	7	8	4	2	4	10	12	8	6
Position	12	2	9	3	2	6	2	10	1	3	5

# Hiding the Repetition Pattern

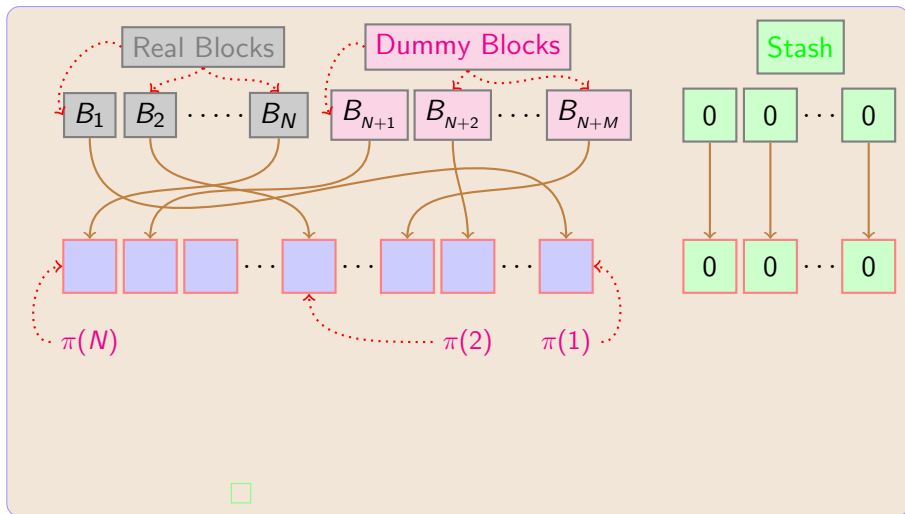
## Initialization for $N$ blocks

- ①  $N$  real blocks  $B_1, \dots, B_N$ ;
- ② create  $M$  dummy blocks  $B_{N+1}, \dots, B_{N+M}$ ;
- ③ create  $M$  stash blocks  $S_1, \dots, S_M$  initialized to 0;
- ④ pick a random permutation  $\pi$  over  $[N + M]$ ;
- ⑤ permute real and dummy blocks according to permutation  $\pi$ 
  - ▶ an encryption of  $B_i$  is uploaded in position  $\pi(i)$ ;
- ⑥ upload all stash blocks in encrypted form;
- ⑦ initialize  $\text{nxt} = 1, \text{cnt} = 1$ ;
- ⑧  $\pi$  is kept private;

# Initial Configuration

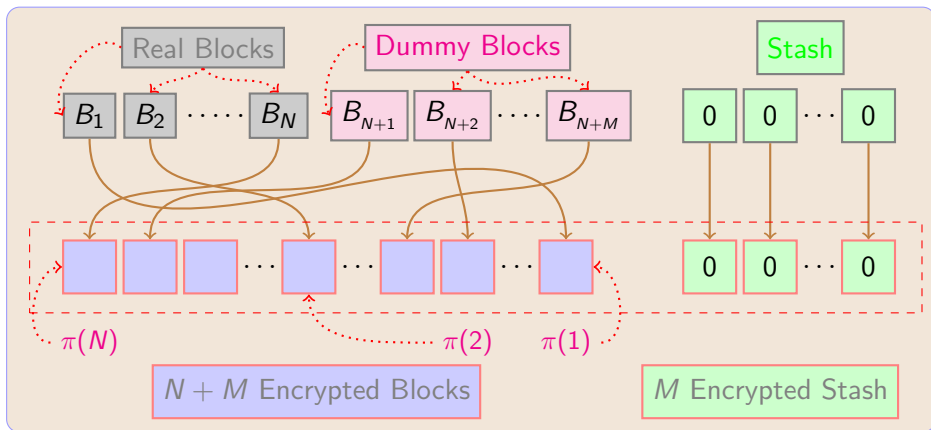


# Initial Configuration

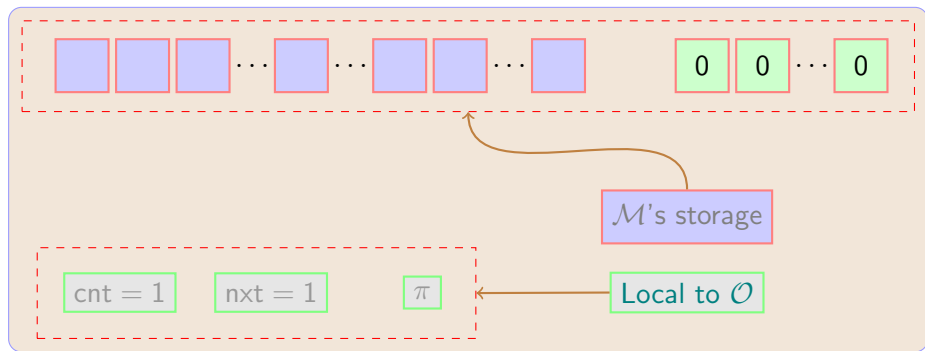




# Initial Configuration



# Initial Configuration

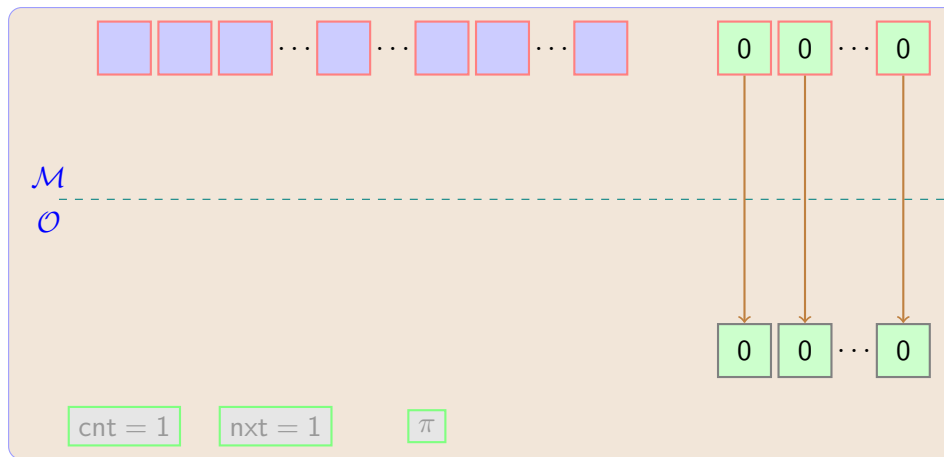


## Reading Block $B_i$

- ① download and decrypt all  $M$  blocks in the Stash;
- ② if  $B_i$  is found in the Stash then
  - ▶ download dummy block  $\pi(N + \text{cnt})$ ;
  - ▶ set  $\text{cnt} = \text{cnt} + 1$ ;else
  - ▶ download encrypted real block in position  $\pi(i)$ ;
  - ▶ decrypt and obtain real block  $B_i$ ;
  - ▶ set next available Stash block  $S_{\text{nxt}} = B_i$ ;
  - ▶ set  $\text{nxt} = \text{nxt} + 1$ ;
- ③ re-encrypt and upload all blocks in the Stash;

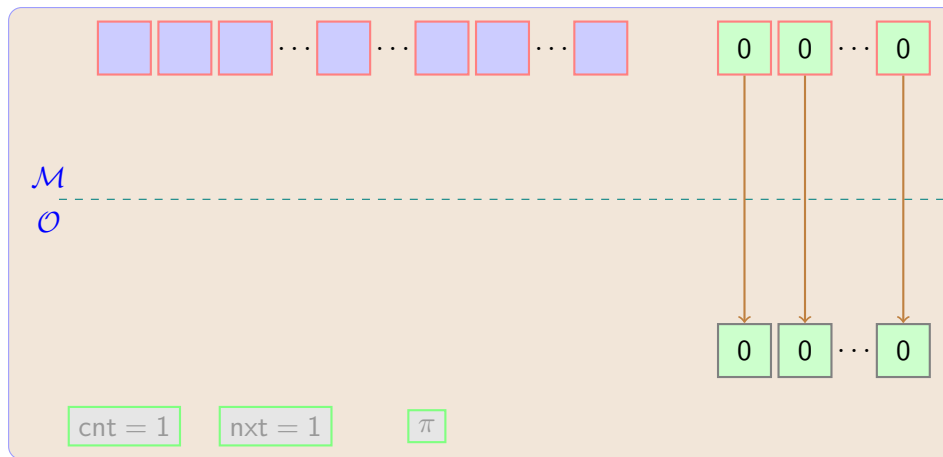
# Reading Block $B_1$

Download and decrypt all blocks from Stash



# Reading Block $B_1$

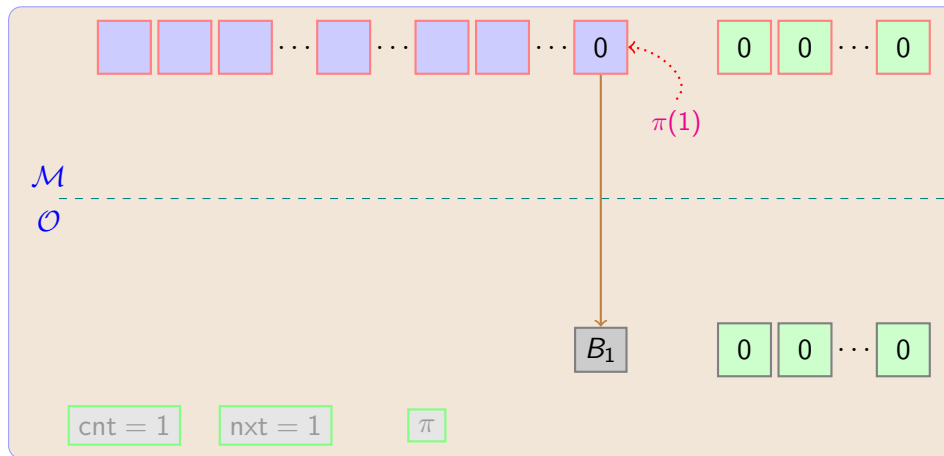
Download and decrypt all blocks from Stash



$B_1$  is not found in the stash

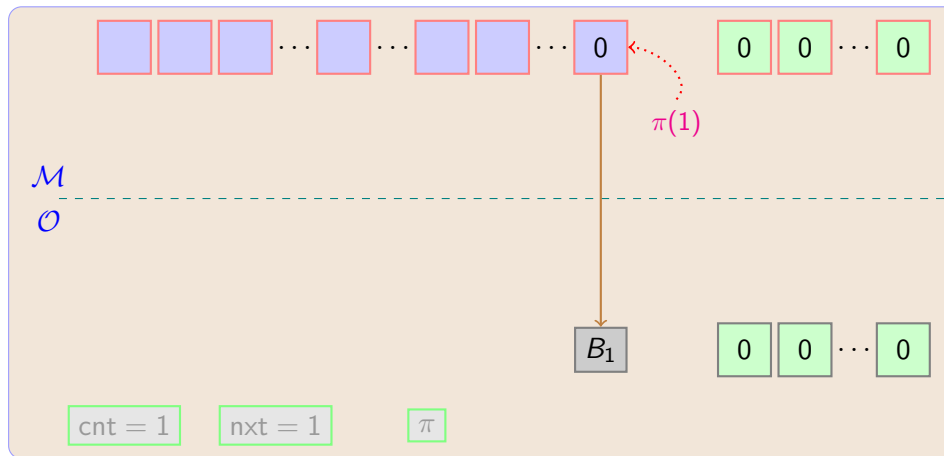
# Reading Block $B_1$

Download block in position  $\pi(1)$



# Reading Block $B_1$

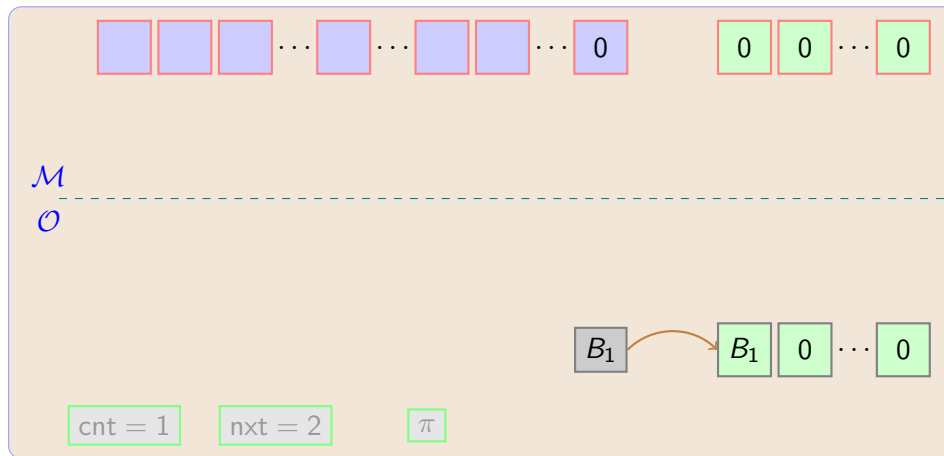
Download block in position  $\pi(1)$



Decrypt and obtain  $B_1$

# Reading Block $B_1$

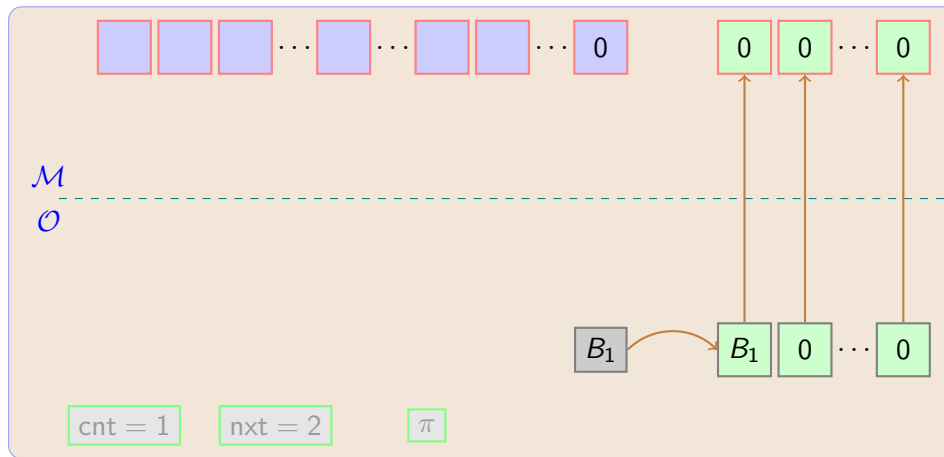
Copy  $B_1$  in the Stash at position  $\text{next}$





# Reading Block $B_1$

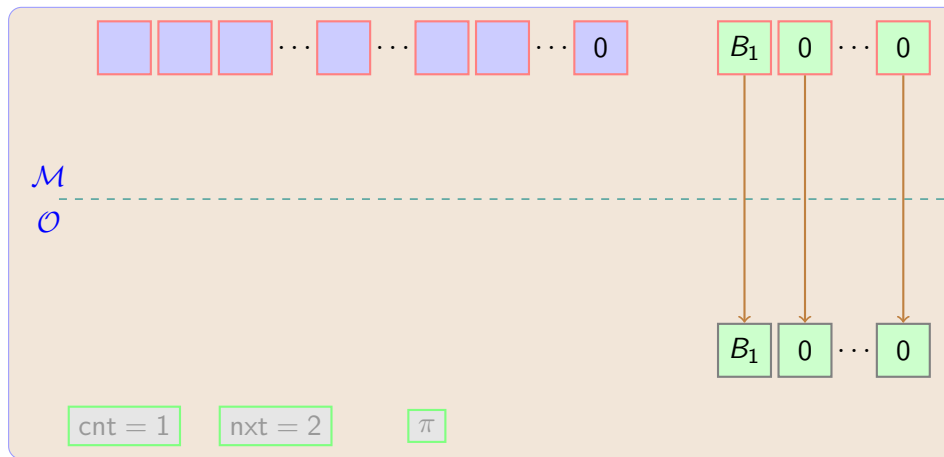
Copy  $B_1$  in the Stash at position  $\text{next}$



## Encrypt and Upload the Stash

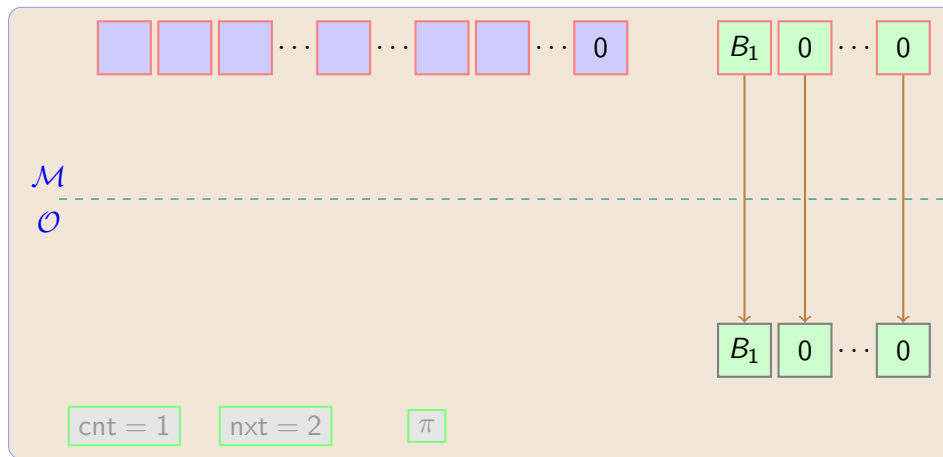
## Reading Block $B_2$

Download and decrypt all blocks from Stash



## Reading Block $B_2$

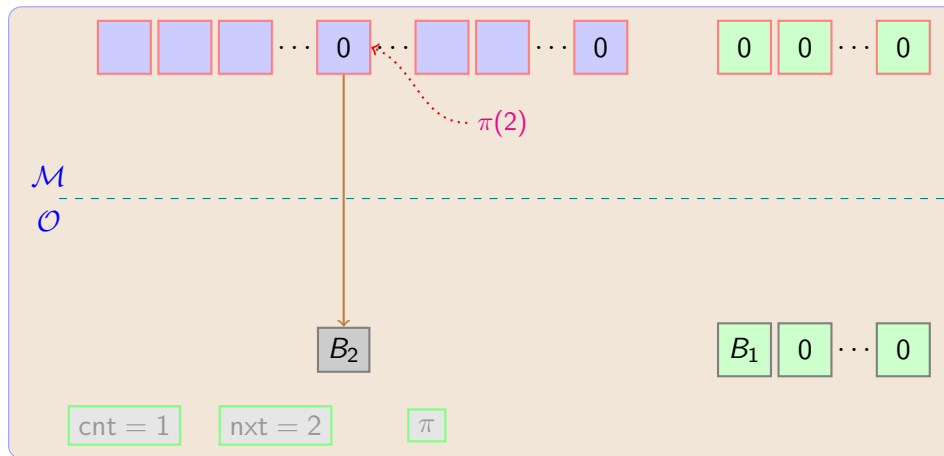
Download and decrypt all blocks from Stash



$B_2$  is not found in the Stash

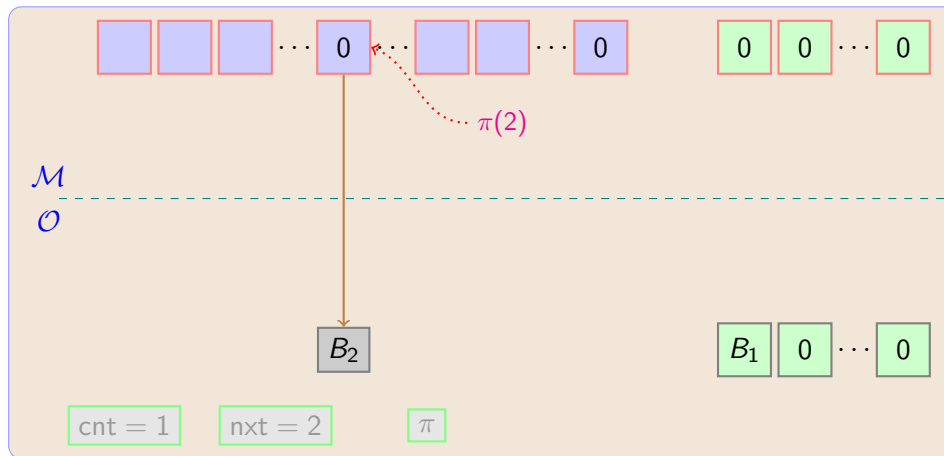
## Reading Block $B_2$

Download block in position  $\pi(2)$



## Reading Block $B_2$

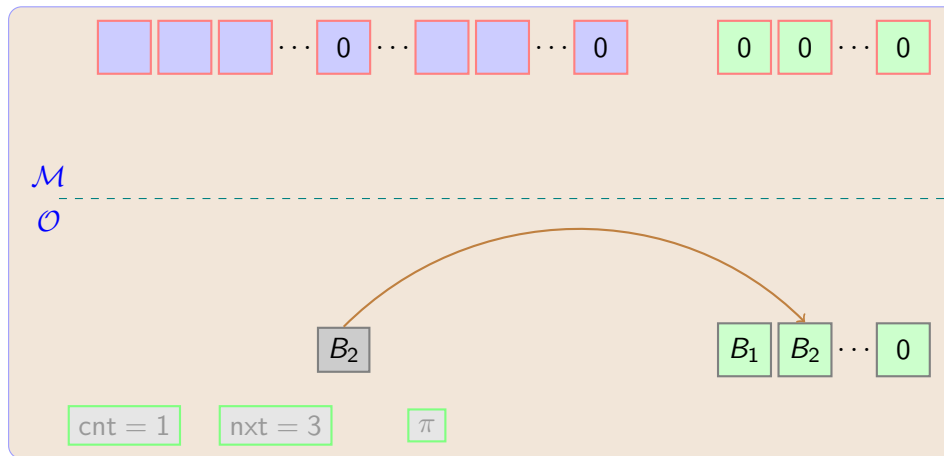
Download block in position  $\pi(2)$



Decrypt and obtain  $B_2$

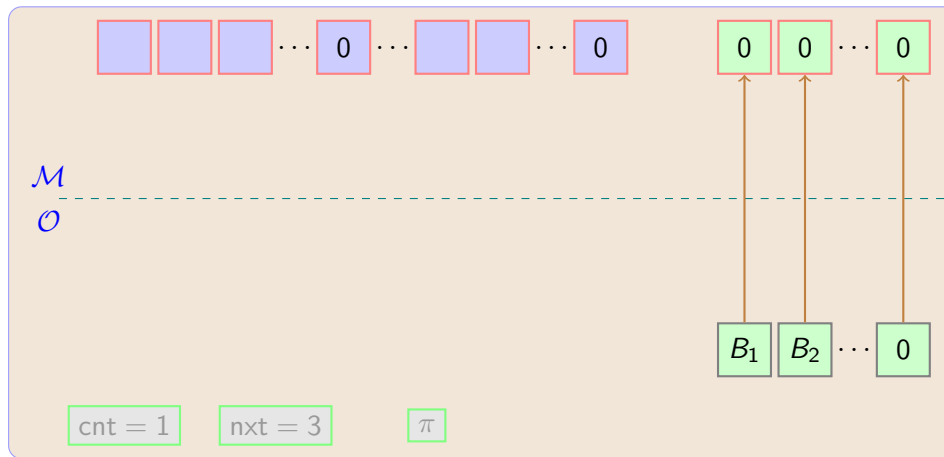
## Reading Block $B_2$

Copy  $B_2$  in the Stash at position  $\text{next}$



## Reading Block $B_2$

Copy  $B_2$  in the Stash at position  $\text{next}$



## Encrypt and Upload the Stash

## Status after reading $B_1$ and $B_2$



$\mathcal{M}$   
—  
 $\mathcal{O}$

cnt = 1

nxt = 2

$\pi$



# Status after reading $B_1$ and $B_2$

Now read  $B_1$  again



$\mathcal{M}$   
—  
 $\mathcal{O}$

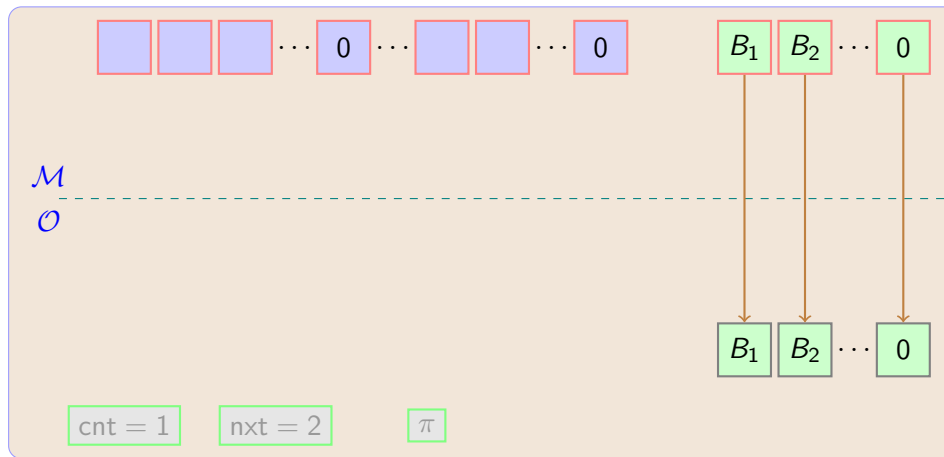
cnt = 1

nxt = 2

$\pi$

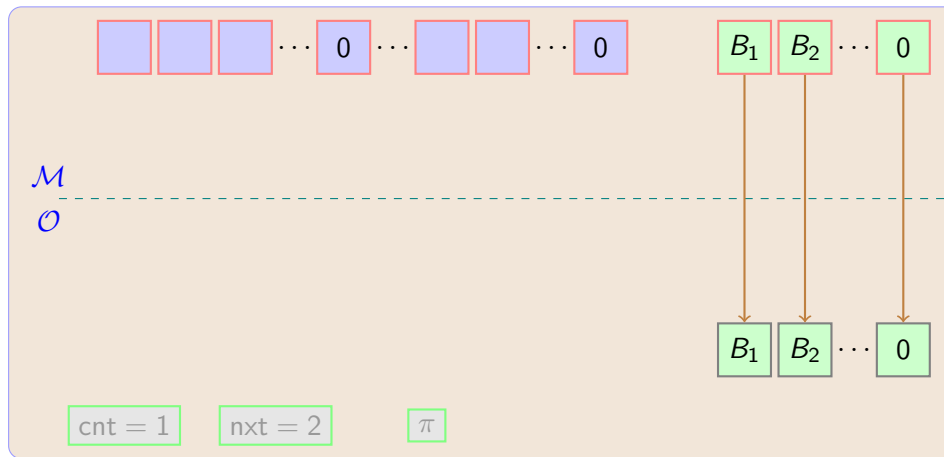
# Status after reading $B_1$ and $B_2$

Download and decrypt all blocks from Stash



## Status after reading $B_1$ and $B_2$

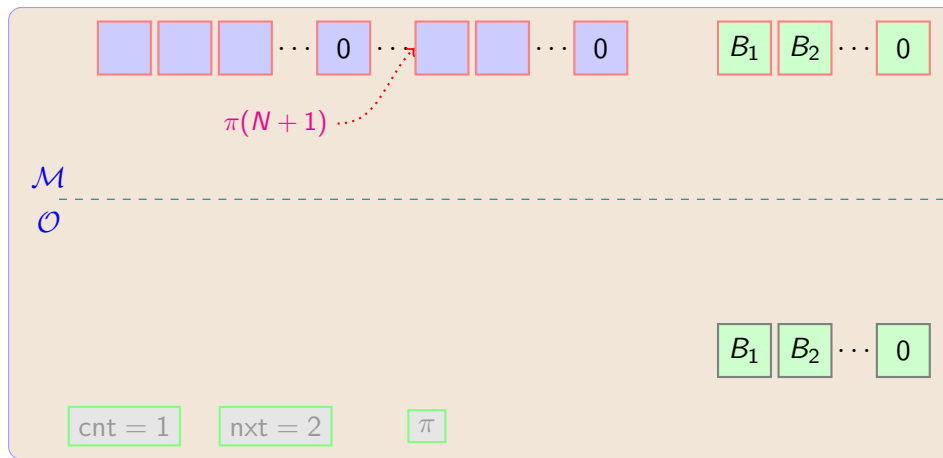
Download and decrypt all blocks from Stash



$B_1$  is found in the Stash

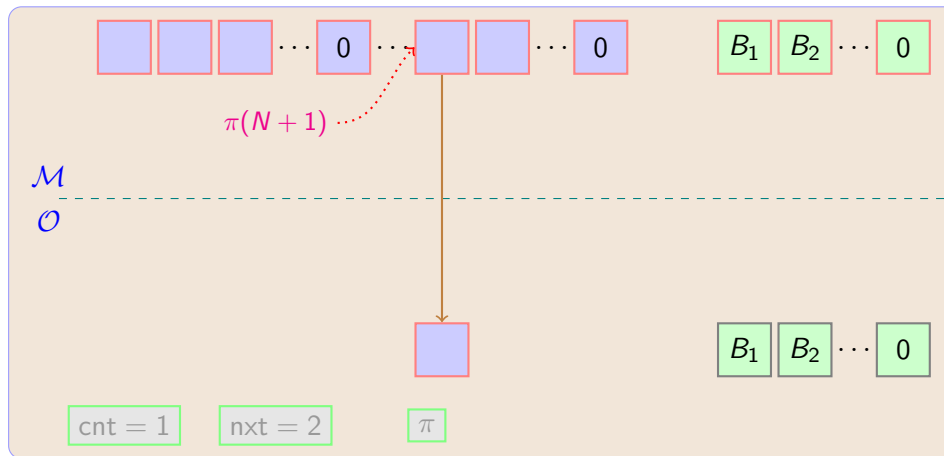
# Reading Block $B_1$ (again)

Download block in position  $\pi(N + \text{cnt})$



## Reading Block $B_1$ (again)

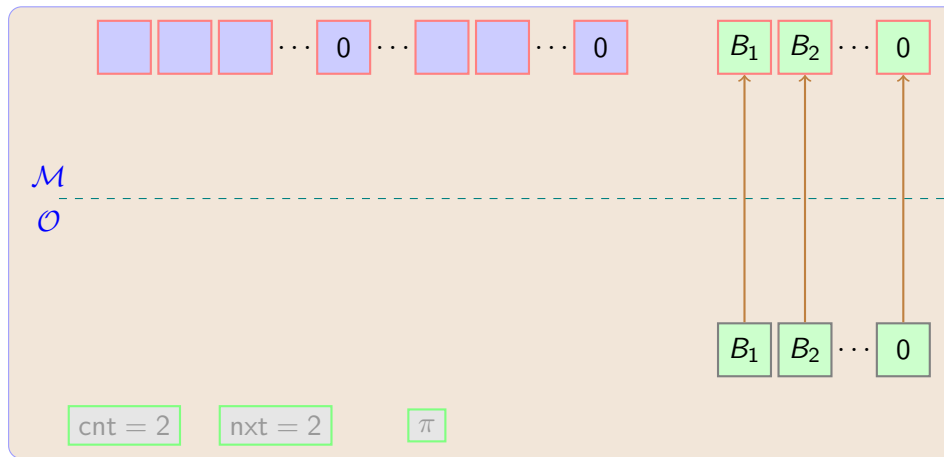
Download block in position  $\pi(N + \text{cnt})$



No need to decrypt

## Reading Block $B_1$ (again)

Download block in position  $\pi(N + \text{cnt})$



## Encrypt and Upload Stash

Insert slide in which we argue obviousness

## Two issues to be dealt with

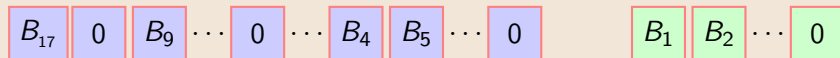
- What happens when the **Stash** is full?



## Two issues to be dealt with

- What happens when the **Stash** is full?
- How much memory does  $\mathcal{O}$  need?
  - ▶ needs to store **cnt** and **nxt**:  $\Theta(1)$  memory;
  - ▶  $\pi$  needs  $O(N)$  memory.

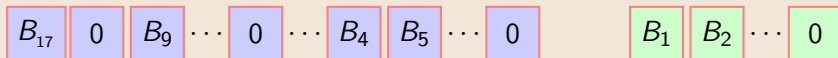
# Overflowing the Stash



$\mathcal{M}$   
—  
 $\mathcal{O}$



# Overflowing the Stash

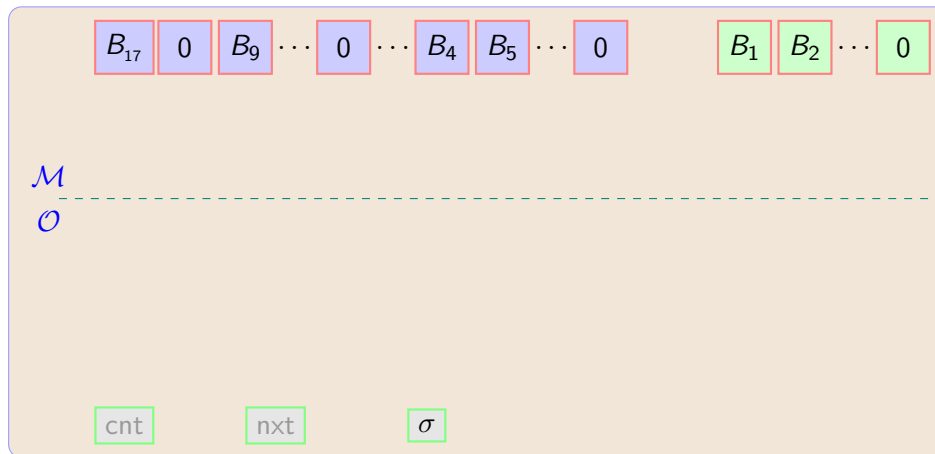


$\mathcal{M}$   
 $\mathcal{O}$

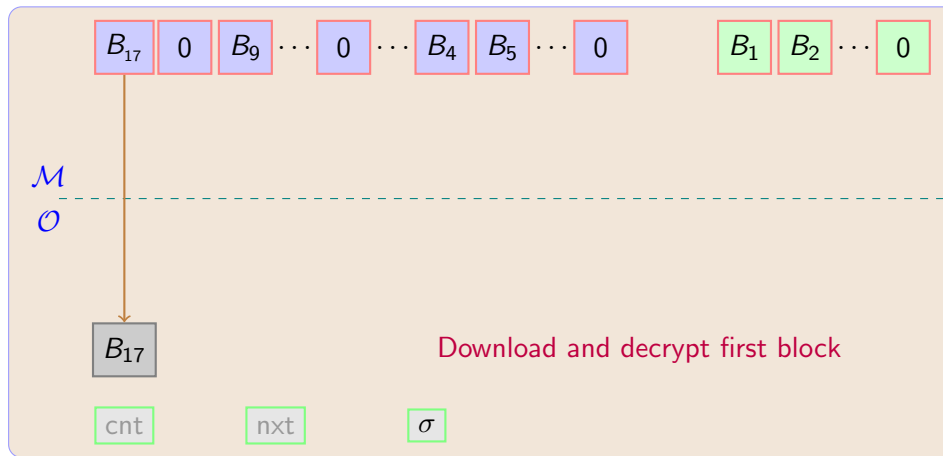
Randomly select a new permutation  $\sigma$



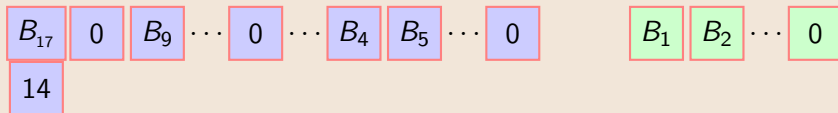
# Overflowing the Stash



# Overflowing the Stash



# Overflowing the Stash

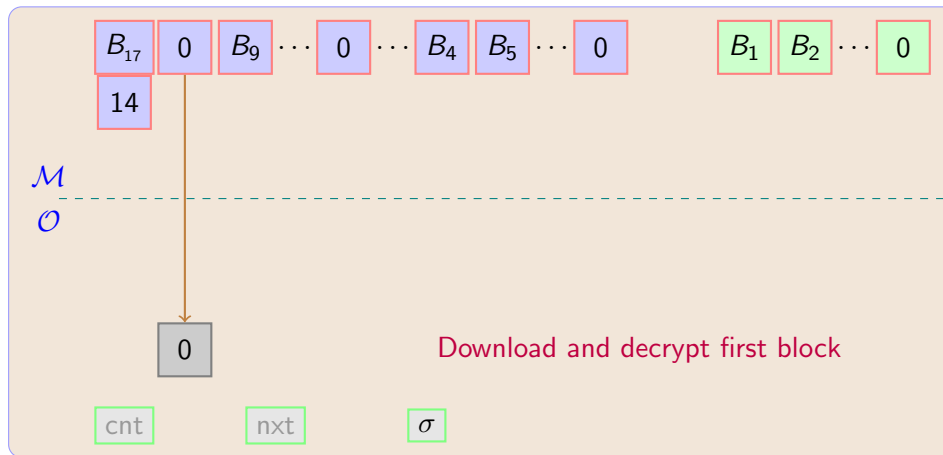


$\mathcal{M}$   
 $\mathcal{O}$

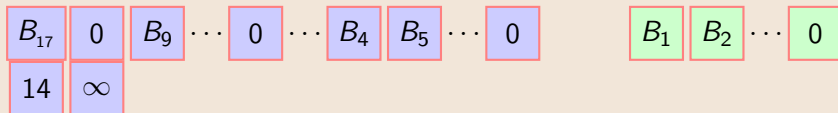
Tag it with  $\sigma(17)$  encrypt and upload



# Overflowing the Stash



# Overflowing the Stash



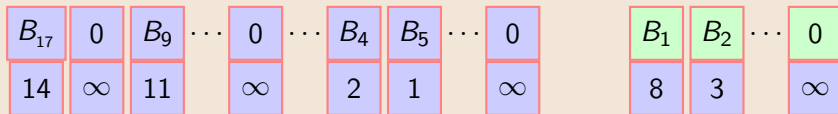
$\mathcal{M}$   
 $\mathcal{O}$

Tag it with  $\infty$  encrypt and upload





# Overflowing the Stash

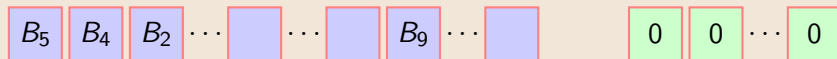


$\mathcal{M}$   
—  
 $\mathcal{O}$

Obliviously sort according to tags

cnt      nxt       $\sigma$

# Overflowing the Stash



$\mathcal{M}$   
—  
 $\mathcal{O}$

cnt = 1

nxt = 1

$\sigma$

# Amortized cost per read operation

Let us count:

# Amortized cost per read operation

Let us count:

- each read costs

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;



# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

for an amortized cost of

$$\Theta \left( M + \frac{N \log N}{M} \right)$$

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

for an amortized cost of

$$\Theta \left( M + \frac{N \log N}{M} \right)$$

for  $M = \sqrt{N}$  we have

$$\sqrt{N} \cdot \log N.$$

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

for an amortized cost of

$$\Theta \left( M + \frac{N \log N}{M} \right)$$

for  $M = \sqrt{N}$  we have

$$\sqrt{N} \cdot \log N.$$

Using AKS to sort.

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

for an amortized cost of

$$\Theta \left( M + \frac{N \log N}{M} \right)$$

for  $M = \sqrt{N}$  we have

$$\sqrt{N} \cdot \log N.$$

Using AKS to sort.

Huge constant

# Amortized cost per read operation

Let us count:

- each read costs
  - ▶  $\Theta(M)$  blocks of bandwidth for the stash;
  - ▶  $\Theta(1)$  blocks of bandwidth for real/dummy blocks;
- after  $M$  reads, we shuffle
  - ▶  $\Theta(M + N) = \Theta(N)$  blocks of bandwidth for tagging;
  - ▶  $\Theta((M + N) \log(M + N)) = \Theta(N \log N)$  blocks of bandwidth for sorting;

for an amortized cost of

$$\Theta \left( M + \frac{N \log N}{M} \right)$$

for  $M = \sqrt{N}$  we have

$$\sqrt{N} \cdot \log N.$$

Using AKS to sort.

In practice  $\sqrt{N} \cdot \log^2 N$ .

Huge constant

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash



# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

using Batcher's sort

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

using Batcher's sort

- Online cost

$$2 \cdot 10^3 \approx 2000$$

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

using Batchers's sort

- Online cost

$$2 \cdot 10^3 \approx 2000$$

- $\mathcal{O}$ 's storage

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

using Batchers's sort

- Online cost

$$2 \cdot 10^3 \approx 2000$$

- $\mathcal{O}$ 's storage
  - ▶ cnt and nxt use constant storage

# In practice...

▶ Jump ahead

## One possible setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks.
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

using Batcher's sort

- Online cost

$$2 \cdot 10^3 \approx 2000$$

- $\mathcal{O}$ 's storage

- ▶ cnt and nxt use constant storage
- ▶  $\pi$  requires storing  $10^6$  4 bytes integers=4 Megabytes

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash



# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

- Online cost: 2 blocks per read

# Keep the stash in $\mathcal{O}$ 's memory

► Jump ahead

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage

# Keep the stash in $\mathcal{O}$ 's memory

► Jump ahead

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - `cnt` and `nxt` use constant storage

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶  $\pi$  requires storing  $10^6$  4-byte integers=4 Megabytes

# Keep the stash in $\mathcal{O}$ 's memory

▶ Jump ahead

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$1/2 \cdot 6^2 \cdot 10^3 \approx 18000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶  $\pi$  requires storing  $10^6$  4-byte integers=4 Megabytes
  - ▶ 1000 blocks of stash for a total of 4 Megabytes



# Shuffling without Sorting

- **Input:**  $N$  blocks stored in  $S[1, \dots, N]$  according to  $\pi$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

Blocks in  $S$  according to  $\pi$

# Shuffling without Sorting

- **Input:**  $N$  blocks stored in  $S[1, \dots, N]$  according to  $\pi$ 
  - ▶ Block  $B_I$  is found in  $[\pi(I)]$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

Blocks in  $S$  according to  $\pi$

# Shuffling without Sorting

- **Input:**  $N$  blocks stored in  $S[1, \dots, N]$  according to  $\pi$ 
  - ▶ Block  $B_i$  is found in  $[\pi(i)]$
- **Output:**  $N$  blocks stored in  $D[1, \dots, N]$  according to  $\sigma$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

Blocks in  $S$  according to  $\pi$

# Shuffling without Sorting

- **Input:**  $N$  blocks stored in  $S[1, \dots, N]$  according to  $\pi$ 
  - ▶ Block  $B_I$  is found in  $[\pi(I)]$
- **Output:**  $N$  blocks stored in  $D[1, \dots, N]$  according to  $\sigma$ 
  - ▶ Block  $B_I$  will be in  $[\sigma(I)]$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

Blocks in  $S$  according to  $\pi$

16	15	5	1
14	2	4	9
10	7	11	6
13	12	8	3

Blocks in  $D$  according to  $\sigma$

# Shuffling $N = 16$

An easy case:

Partition the  $N$  blocks in  $\sqrt{N}$  groups of  $\sqrt{N}$

Blocks in  $S$  according to  $\pi$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

$\mathcal{M}$

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Partition the  $N$  destinations in  $\sqrt{N}$  groups of  $\sqrt{N}$

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download first source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download first source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$\mathcal{O}$

3	4	2	10
---	---	---	----



# Shuffling $N = 16$

An easy case:

One block to each destination group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
0	0	0	0
10	2	4	3

$\mathcal{O}$

9	8	15	16
---	---	----	----

# Shuffling $N = 16$

An easy case:

One block to each destination group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
0	0	0	0
16	15	8	9
10	2	4	3

$\mathcal{O}$

1	5	7	14
---	---	---	----

# Shuffling $N = 16$

An easy case:

One block to each destination group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$



# Shuffling $N = 16$

An easy case:

Download second source group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

0	0	0	0
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

6	11	12	13
---	----	----	----

# Shuffling $N = 16$

An easy case:

One block to each destination group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

13	12	11	6
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Each block in the right destination group

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

13	12	11	6
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

13	12	11	6
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

13	12	11	6
14	7	5	1
16	15	8	9
10	2	4	3

$\mathcal{O}$

10	16	14	13
----	----	----	----

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	12	11	6
14	7	5	1
10	15	8	9
13	2	4	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	12	11	6
14	7	5	1
10	15	8	9
13	2	4	3

$\mathcal{O}$

2	15	7	12
---	----	---	----

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	15	11	6
14	2	5	1
10	7	8	9
13	12	4	3

$\mathcal{O}$

4	8	5	11
---	---	---	----



# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	15	5	6
14	2	4	1
10	7	11	9
13	12	8	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	15	5	6
14	2	4	1
10	7	11	9
13	12	8	3

$\mathcal{O}$

3	9	1	6
---	---	---	---

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

$\mathcal{M}$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

16	15	5	1
14	2	4	9
10	7	11	6
13	12	8	3

$\mathcal{O}$

# Shuffling $N = 16$

An easy case:

Download each group and upload in correct position

Blocks in  $S$  according to  $\pi$

10	16	14	13
2	15	7	12
4	8	5	11
3	9	1	6

Blocks in  $D$  according to  $\sigma$

16	15	5	1
14	2	4	9
10	7	11	6
13	12	8	3

$\mathcal{M}$

$\mathcal{O}$

# Analysis of Shuffling Algorithm

- **Obliviousness:** access pattern independent of  $\pi, \sigma$

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group

# Analysis of Shuffling Algorithm

- **Obliviousness:** access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location

# Analysis of Shuffling Algorithm

- **Obliviousness:** access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group



# Analysis of Shuffling Algorithm

- **Obliviousness:** access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group
- **bandwidth**:  $4N$   
each block

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group
- **bandwidth**:  $4N$   
each block
  - ▶ downloaded exactly twice

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group
- **bandwidth**:  $4N$   
each block
  - ▶ downloaded exactly twice
  - ▶ uploaded exactly twice

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group
- **bandwidth**:  $4N$   
each block
  - ▶ downloaded exactly twice
  - ▶ uploaded exactly twice
- **Luck**: so much!!!

# Analysis of Shuffling Algorithm

- **Obliviousness**: access pattern independent of  $\pi, \sigma$ 
  - ▶ download each source group
  - ▶ upload one block to each destination group in the next available empty location
  - ▶ download each destination group
  - ▶ upload each destination group
- **bandwidth**:  $4N$   
each block
  - ▶ downloaded exactly twice
  - ▶ uploaded exactly twice
- **Luck**: so much!!!
  - ▶ each source group contains exactly one block for each destination group under  $\sigma$

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups



# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache
- Recalibrate phase

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache
- Recalibrate phase
  - ▶ download each destination group

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache
- Recalibrate phase
  - ▶ download each destination group
  - ▶ remove dummies



# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache
- Recalibrate phase
  - ▶ download each destination group
  - ▶ remove dummies
  - ▶ add blocks found in local cache

# Efficient Shuffling: CacheShuffleRoot

*when you know you are not going to be lucky, just randomize*

- Randomly partition destination array in  $(1 + 2\epsilon)\sqrt{N}$  groups
  - ▶ Each group has size at most  $(1 - \epsilon)\sqrt{N}$ , except with negligible probability
- Spray phase
  - ▶ Download each source group and spray exactly one block to each destination group
  - ▶ if none available, spray a dummy block
  - ▶ if more than one available, spray exactly one and store the extra blocks in local cache
- Recalibrate phase
  - ▶ download each destination group
  - ▶ remove dummies
  - ▶ add blocks found in local cache
  - ▶ upload in correct order

# Analysis of CacheShuffleRoot

- bandwidth:  $(4 + \epsilon)N$  blocks

# Analysis of CacheShuffleRoot

- **bandwidth**:  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group



# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group
  - ▶ upload one block to each destination group

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group
  - ▶ upload one block to each destination group
  - ▶ when done, download and upload each destination group

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group
  - ▶ upload one block to each destination group
  - ▶ when done, download and upload each destination group
- **$\mathcal{M}$ 's storage:**  $\Theta(N)$

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group
  - ▶ upload one block to each destination group
  - ▶ when done, download and upload each destination group
- **$\mathcal{M}$ 's storage:**  $\Theta(N)$
- **$\mathcal{O}$ 's storage**

# Analysis of CacheShuffleRoot

- **bandwidth:**  $(4 + \epsilon)N$  blocks
  - ▶ each **real** block is downloaded exactly twice and uploaded exactly twice
  - ▶ each **dummy** block is uploaded exactly once and downloaded exactly once
  - ▶ number of **dummy** blocks is at most  $\epsilon \cdot N$  except with negligible probability
- **Obliviousness:** access pattern is fixed for all executions
  - ▶ download each source group
  - ▶ upload one block to each destination group
  - ▶ when done, download and upload each destination group
- **$\mathcal{M}$ 's storage:**  $\Theta(N)$
- **$\mathcal{O}$ 's storage**
  - ▶ next slide

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1



# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1
- each cache is expected to hold a constant number of blocks

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1
- each cache is expected to hold a constant number of blocks
- total  $\mathcal{O}$ 's storage is  $\Theta(\sqrt{N})$  except with negligible probability

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1
- each cache is expected to hold a constant number of blocks
- total  $\mathcal{O}$ 's storage is  $\Theta(\sqrt{N})$  except with negligible probability

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1
- each cache is expected to hold a constant number of blocks
- total  $\mathcal{O}$ 's storage is  $\Theta(\sqrt{N})$  except with negligible probability

## Formal argument

- cache sizes are not independent so cannot use Chernoff bound

# Analysis of CacheShuffleRoot: $\mathcal{O}$ 's storage

## Intuition

- $\mathcal{O}$ 's keeps a cache for each of the  $(1 + 2\epsilon)\sqrt{N}$  destination groups
- each cache has arrival rate  $\approx (1 - 2\epsilon)$
- departure rate 1
- each cache is expected to hold a constant number of blocks
- total  $\mathcal{O}$ 's storage is  $\Theta(\sqrt{N})$  except with negligible probability

## Formal argument

- cache sizes are not independent so cannot use Chernoff bound
- prove negative association and then use Chernoff bound

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks



# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

- Online cost: 2 blocks per read

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶  $\pi$  requires storing  $10^6$  4-byte integers=4 Megabytes

# Keep the stash in $\mathcal{O}$ 's memory and use CacheShuffleRoot

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash

## Resources needed:

- $\mathcal{M}$ 's storage:  $N + M = 10^6 + 10^3$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 4000$$

- Online cost: 2 blocks per read
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶  $\pi$  requires storing  $10^6$  4-byte integers=4 Megabytes
  - ▶ 1000 blocks of stash for a total of 4 Megabytes

# Where are we now?

## ► Construction -1

### Keep It

- $\mathcal{M}$ 's storage: 0
- $\mathcal{O}$ 's storage:  $N$
- bandwidth 0

## ► Construction 0

### Download It

- $\mathcal{M}$ 's storage:  $N$
- $\mathcal{O}$ 's storage: 1
- bandwidth  $N$

## ► Construction 1

### Download Stash

- $\mathcal{M}$ 's storage:  $N + \sqrt{N}$
- $\mathcal{O}$ 's storage: 1
- Online Comm.  $O(\sqrt{N})$
- Am. Comm.  
 $O(\sqrt{N} \cdot \log N)$

## ► Construction 2

### Keep Stash


- $\mathcal{M}$ 's storage:  $N + \sqrt{N}$
- $\mathcal{O}$ 's storage:  $\sqrt{N}$
- Online Comm. 1
- Am. Comm.  
 $O(\sqrt{N} \cdot \log N)$

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern



# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does  hide access to the **Stash**?

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does **0** hide access to the **Stash**?
  - ▶ Use an ORAM!

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does **0** hide access to the **Stash**?
  - ▶ Use an ORAM!
  - ▶ We only have two possible ORAMs:

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does **🕒** hide access to the **Stash**?
  - ▶ Use an ORAM!
  - ▶ We only have two possible ORAMs:
  - ▶ **Download Stash** uses the **Download It** ORAM to manage the **Stash**

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does **🕒** hide access to the **Stash**?
  - ▶ Use an ORAM!
  - ▶ We only have two possible ORAMs:
  - ▶ **Download Stash** uses the **Download It** ORAM to manage the **Stash**
  - ▶ **Keep Stash** uses the **Keep It** ORAM to manage the **Stash**

# Download and Keep Stash

- Use a cache (the **Stash**) to hide the repetition pattern
- How does **🕒** hide access to the **Stash**?
  - ▶ Use an ORAM!
  - ▶ We only have two possible ORAMs:
  - ▶ **Download Stash** uses the **Download It** ORAM to manage the **Stash**
  - ▶ **Keep Stash** uses the **Keep It** ORAM to manage the **Stash**

# Download and Keep Stash

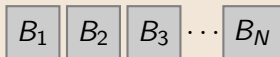
- Use a cache (the **Stash**) to hide the repetition pattern
- How does **🕒** hide access to the **Stash**?
  - ▶ Use an ORAM!
  - ▶ We only have two possible ORAMs:
  - ▶ **Download Stash** uses the **Download It** ORAM to manage the **Stash**
  - ▶ **Keep Stash** uses the **Keep It** ORAM to manage the **Stash**

But now we have more ORAMs!!!

$\mathcal{M}$

$\mathcal{O}$

$N$  blocks of data







$$N + \rho N$$

$\mathcal{M}$

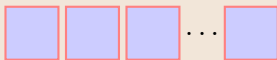
$\mathcal{O}$

$\rho N$  blocks of stash





$$N + \rho N$$



$$\rho N + \rho^2 N$$

$\mathcal{M}$

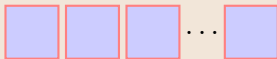
$\mathcal{O}$

$\rho^2 N$  blocks of stash





$$N + \rho N$$



$$\rho N + \rho^2 N$$



$$\rho^2 N + \rho^3 N$$

$\mathcal{M}$

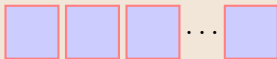
$\mathcal{O}$

$\rho^3 N$  blocks of stash





$$N + \rho N$$



$$\rho N + \rho^2 N$$



$$\rho^2 N + \rho^3 N$$

$\mathcal{M}$

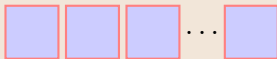
$\mathcal{O}$

$\rho^3 N$  blocks of stash





$$N + \rho N$$



$$\rho N + \rho^2 N$$



$$\rho^2 N + \rho^3 N$$

$\mathcal{M}$

$\mathcal{O}$

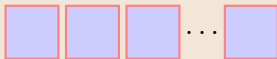
$\sqrt{N}$  blocks of stash

$$\rho = N^{-1/6}$$





$$N + N^{5/6}$$



$$N^{5/6} + N^{2/3}$$



$$N^{2/3} + N^{1/2}$$

$\mathcal{M}$

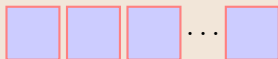
$\mathcal{O}$

$\sqrt{N}$  blocks of stash





Level 3  $N + N^{5/6}$



Level 2  $N^{5/6} + N^{2/3}$



Level 1  $N^{2/3} + N^{1/2}$

$\mathcal{M}$

$\mathcal{O}$

$\sqrt{N}$  blocks of stash



Level 0

Position Map

$(\text{lev}_i, \text{pos}_i) \ i = 1, \dots, N$

## 3-level ORAM: Querying

Querying  $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;



## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$

## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next dummy in level  $l$ ;

## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ asks for block in  $\text{pos}_q$  from level  $\text{lev}_q$ ;

## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ asks for block in  $\text{pos}_q$  from level  $\text{lev}_q$ ;
  - ▶  $B_q$  is then stored in level 0 and  $(\text{lev}_q, \text{pos}_q)$  is updated;

## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ asks for block in  $\text{pos}_q$  from level  $\text{lev}_q$ ;
  - ▶  $B_q$  is then stored in level 0 and  $(\text{lev}_q, \text{pos}_q)$  is updated;
- else

## 3-level ORAM: Querying

### Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ asks for block in  $\text{pos}_q$  from level  $\text{lev}_q$ ;
  - ▶  $B_q$  is then stored in level 0 and  $(\text{lev}_q, \text{pos}_q)$  is updated;
- else
  - ▶ for  $l = 1, 2, 3$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;

# 3-level ORAM: Querying

## Querying $B_q$

- retrieve  $(\text{lev}_q, \text{pos}_q)$  from local memory;
- if  $\text{lev}_q \neq 0$ 
  - ▶ for all  $l \neq \text{lev}_q$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ asks for block in  $\text{pos}_q$  from level  $\text{lev}_q$ ;
  - ▶  $B_q$  is then stored in level 0 and  $(\text{lev}_q, \text{pos}_q)$  is updated;
- else
  - ▶ for  $l = 1, 2, 3$ ,  $\mathcal{O}$  asks for the next **dummy** in level  $l$ ;
  - ▶ block  $B_q$  is retrieved from local stash (level 0);

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;



## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries

## 3-level ORAM: Analysis

- at the start only level 3 contains real blocks;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;

## 3-level ORAM: Analysis

- at the start only level 3 contains real blocks;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 1 is full after  $N^{2/3}$  queries

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;

## 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$



### 3-level ORAM: Analysis

- at the start only level 3 contains real blocks;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 1 is full after  $N^{2/3}$  queries
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times

### 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 1 is full after  $N^{2/3}$  queries
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 2 is full after  $N^{5/6}$  queries

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 2 is full after  $N^{5/6}$  queries**
  - ▶ it is shuffled with the level 3 of size  $N$ ;

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 2 is full after  $N^{5/6}$  queries**
  - ▶ it is shuffled with the level 3 of size  $N$ ;
  - ▶ each shuffle costs  $4N$

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 2 is full after  $N^{5/6}$  queries**
  - ▶ it is shuffled with the level 3 of size  $N$ ;
  - ▶ each shuffle costs  $4N$
  - ▶ over  $N$  queries, it happens  $N^{1/6}$  times

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- **the local stash is full after  $N^{1/2}$  queries**
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 1 is full after  $N^{2/3}$  queries**
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- **level 2 is full after  $N^{5/6}$  queries**
  - ▶ it is shuffled with the level 3 of size  $N$ ;
  - ▶ each shuffle costs  $4N$
  - ▶ over  $N$  queries, it happens  $N^{1/6}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$

# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 1 is full after  $N^{2/3}$  queries
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 2 is full after  $N^{5/6}$  queries
  - ▶ it is shuffled with the level 3 of size  $N$ ;
  - ▶ each shuffle costs  $4N$
  - ▶ over  $N$  queries, it happens  $N^{1/6}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- Over  $N$  queries, the cost is  $12 \cdot N^{7/6}$



# 3-level ORAM: Analysis

- at the start only level 3 contains **real blocks**;
- the local stash is full after  $N^{1/2}$  queries
  - ▶ it is shuffled with the level 1 of size  $N^{2/3}$ ;
  - ▶ each shuffle costs  $4N^{2/3}$
  - ▶ over  $N$  queries, it happens  $N^{1/2}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 1 is full after  $N^{2/3}$  queries
  - ▶ it is shuffled with the level 2 of size  $N^{5/6}$ ;
  - ▶ each shuffle costs  $4N^{5/6}$
  - ▶ over  $N$  queries, it happens  $N^{1/3}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- level 2 is full after  $N^{5/6}$  queries
  - ▶ it is shuffled with the level 3 of size  $N$ ;
  - ▶ each shuffle costs  $4N$
  - ▶ over  $N$  queries, it happens  $N^{1/6}$  times
  - ▶ total cost:  $4 \cdot N^{7/6}$
- Over  $N$  queries, the cost is  $12 \cdot N^{7/6}$ 
  - ▶ each query has an amortized cost of  $12N^{1/6}$  blocks;

## 3-level ORAM: in practice

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

## 3-level ORAM: in practice

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

Resources needed:

## 3-level ORAM: in practice

Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks

## 3-level ORAM: in practice

### Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

### Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

## 3-level ORAM: in practice

### Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

### Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

- Online cost: 3 blocks downloaded

## 3-level ORAM: in practice

### Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

### Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

- Online cost: 3 blocks downloaded
- $\mathcal{O}$ 's storage

## 3-level ORAM: in practice

### Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

### Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

- Online cost: 3 blocks downloaded
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage



# 3-level ORAM: in practice

## Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

## Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

- Online cost: 3 blocks downloaded
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶ `position map` requires storing  $10^6$  6-byte integers  $\approx 4$  Megabytes

## 3-level ORAM: in practice

### Same setting:

- $N = 10^6$  blocks of 4K each for a total of 4 Gigabytes
- $M = 10^3$  blocks of stash for a total of 4 Megabytes

### Resources needed:

- $\mathcal{M}$ 's storage:  $10^6 + 2 \cdot 10^5 + 2 \cdot 10^4 + 10^3 = 1,221,000$  blocks
- Cost of shuffling amortized per read operation:

$$4 \cdot 10^3 \approx 120$$

- Online cost: 3 blocks downloaded
- $\mathcal{O}$ 's storage
  - ▶ `cnt` and `nxt` use constant storage
  - ▶ `position map` requires storing  $10^6$  6-byte integers  $\approx 4$  Megabytes
  - ▶ 1000 blocks of stash for a total of 4 Megabytes

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks



# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks
- Amortized bandwidth  $0.55 \cdot \log_2 N$

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks
- Amortized bandwidth  $0.55 \cdot \log_2 N$ 
  - ▶ For  $N = 10^6$ , 21 Blocks

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks
- Amortized bandwidth  $0.55 \cdot \log_2 N$ 
  - ▶ For  $N = 10^6$ , 21 Blocks
- OnLine bandwidth: 1 Block

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks
- Amortized bandwidth  $0.55 \cdot \log_2 N$ 
  - ▶ For  $N = 10^6$ , 21 Blocks
- OnLine bandwidth: 1 Block

# Why stop at 3 levels?

$\mathcal{O}$ 's memory  $\approx \sqrt{N}$

- set  $\rho = 1/2$
- $1/2 \log_2 N$  levels with
  - ▶  $N + N/2$  blocks
  - ▶  $N/2 + N/4$  blocks
  - ▶ ...
  - ▶  $3\sqrt{N}$  blocks
- Amortized bandwidth  $0.55 \cdot \log_2 N$ 
  - ▶ For  $N = 10^6$ , 21 Blocks
- OnLine bandwidth: 1 Block

## Techniques to reduce bandwidth

- XOR Technique
- Homomorphic Selection
- Compression via Polynomial Interpolation

# The XOR technique to reduce bandwidth

## $\ell$ -level ORAM

- when asking for  $B_q$ ,  $\mathcal{O}$  asks one block for each level

# The XOR technique to reduce bandwidth

## $\ell$ -level ORAM

- when asking for  $B_q$ ,  $\mathcal{O}$  asks one block for each level
- at most one block is real

# The XOR technique to reduce bandwidth

## $l$ -level ORAM

- when asking for  $B_q$ ,  $\mathcal{O}$  asks one block for each level
- at most one block is real
- suppose  $\mathcal{O}$  can compute locally the  $l - 1$  dummy blocks requested



# The XOR technique to reduce bandwidth

## $l$ -level ORAM

- when asking for  $B_q$ ,  $\mathcal{O}$  asks one block for each level
- at most one block is real
- suppose  $\mathcal{O}$  can compute locally the  $l - 1$  dummy blocks requested
  - ▶  $\mathcal{M}$  instead of sending all  $l$  blocks individually, xors them together and sends the result to  $\mathcal{O}$

# The XOR technique to reduce bandwidth

## $l$ -level ORAM

- when asking for  $B_q$ ,  $\mathcal{O}$  asks one block for each level
- at most one block is real
- suppose  $\mathcal{O}$  can compute locally the  $l - 1$  dummy blocks requested
  - ▶  $\mathcal{M}$  instead of sending all  $l$  blocks individually, xors them together and sends the result to  $\mathcal{O}$
  - ▶  $\mathcal{O}$  computes the  $l - 1$  dummy blocks and xors them with the block received from  $\mathcal{M}$

# Assumption:

suppose  $\mathcal{O}$  can compute any *dummy* block without interacting with  $\mathcal{M}$

- each block uniquely identified by  $(l, pos)$
- a *dummy* block is an AES-ECB encryption of  $0^{\text{len}}$
- using key  $\mathcal{F}(K, (l, pos))$ 
  - ▶  $\mathcal{F}$  is a pseudorandom function
  - ▶  $K$  is a randomly chosen seed private to  $\mathcal{O}$

# Some Theory

## A Taxonomy

### • OnLine vs OffLine ORAM

- ▶ In an **OnLine** ORAM, all requests come one at the time and must be satisfied before the next one
- ▶ in an **OffLine** ORAM, all requests come together

### • BallsAndBins

- ▶ Blocks are atomic and opaque blobs of data

### • Passive vs Active $\mathcal{M}$

- ▶ A **Passive**  $\mathcal{M}$  only moves data
- ▶ An **Active**  $\mathcal{M}$  can perform computation on data
  - ★ The **XOR technique** requires an **Active**  $\mathcal{M}$

# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$

# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$ 
  - ▶ **BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$

# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$ 
  - ▶ **BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$
  - ▶ **NonBallsAndBins** and **OnLine** with **Passive**  $\mathcal{M}$

# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$ 
  - ▶ **BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$
  - ▶ **NonBallsAndBins** and **OnLine** with **Passive**  $\mathcal{M}$
- There are  $\tilde{O}(\log N)$  **OnLine** ORAM in the Balls and Bins model with **Passive**  $\mathcal{M}$



# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$ 
  - ▶ **BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$
  - ▶ **NonBallsAndBins** and **OnLine** with **Passive**  $\mathcal{M}$
- There are  $\tilde{O}(\log N)$  **OnLine** ORAM in the Balls and Bins model with **Passive**  $\mathcal{M}$
- There are  $o(\log N)$  **OnLine** ORAM with **Active**  $\mathcal{M}$

# What we know in the asymptotic world

- All the following require  $\Omega(\log N)$ 
  - ▶ **BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$
  - ▶ **NonBallsAndBins** and **OnLine** with **Passive**  $\mathcal{M}$
- There are  $\tilde{O}(\log N)$  **OnLine** ORAM in the Balls and Bins model with **Passive**  $\mathcal{M}$
- There are  $o(\log N)$  **OnLine** ORAM with **Active**  $\mathcal{M}$
- Proving lower bound for **Non-BallsAndBins** and **OffLine** with **Passive**  $\mathcal{M}$  would give a superlinear lower bound for sorting circuits.

## $(\epsilon, \delta)$ -Differential Privacy

- $\mathcal{M}$  stores  $n$  blocks of memory.
- Every time  $\mathcal{O}$  wants a block, he asks  $\mathcal{M}$  one or more blocks.
- Security notion:
  - ▶ For any two block sequences  $\mathbb{B} = B_1, \dots, B_n$  and  $\mathbb{C} = C_1, \dots, C_n$
  - ▶ For any two access sequences  $i_1, \dots, i_l$  and  $j_1, \dots, j_l$  **that differ in one position**
    - ★ performing access  $i_1, \dots, i_l$  on  $\mathbb{B} = B_1, \dots, B_n$ ;
    - ★ performing access  $j_1, \dots, j_l$  on  $\mathbb{C} = C_1, \dots, C_n$

generate the same distribution of accesses to the data stored by  $\mathcal{M}$

## $(\epsilon, \delta)$ -Differential Privacy

- $\mathcal{M}$  stores  $n$  blocks of memory.
- Every time  $\mathcal{O}$  wants a block, he asks  $\mathcal{M}$  one or more blocks.
- Security notion:
  - ▶ For any two block sequences  $\mathbb{B} = B_1, \dots, B_n$  and  $\mathbb{C} = C_1, \dots, C_n$
  - ▶ For any two access sequences  $i_1, \dots, i_l$  and  $j_1, \dots, j_l$  that differ in one position
    - ★ performing access  $i_1, \dots, i_l$  on  $\mathbb{B} = B_1, \dots, B_n$ ;
    - ★ performing access  $j_1, \dots, j_l$  on  $\mathbb{C} = C_1, \dots, C_n$

generate the same distribution of accesses to the data stored by  $\mathcal{M}$

For every predicate  $A$

$$\begin{aligned} & \text{Prob}[\text{view} \leftarrow \text{View}(I, \mathbb{B}) : A(\text{view}) = 1] \\ & \leq e^\epsilon \cdot \text{Prob}[\text{view} \leftarrow \text{View}(J, \mathbb{C}) : A(\text{view}) = 1] + \delta \end{aligned}$$

# Differential Privacy

- it protects only **individual** accesses

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?



# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?
  - ▶ it depends on  $\epsilon$

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?
  - ▶ it depends on  $\epsilon$
- no protection is offered for **high-probability** accesses

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?
  - ▶ it depends on  $\epsilon$
- no protection is offered for **high-probability** accesses
  - ▶ **nothing is lost: my daily routine is public**

# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?
  - ▶ it depends on  $\epsilon$
- no protection is offered for **high-probability** accesses
  - ▶ **nothing is lost: my daily routine is public**
- I really want to protect **unusual accesses** to documents that might reveal something

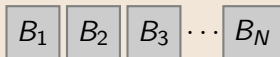
# Differential Privacy

- it protects only **individual** accesses
- if the probability of an access is larger than  $1/e^\epsilon$  no guarantee is given
- it protects **unusual accesses**
- how unusual?
  - ▶ it depends on  $\epsilon$
- no protection is offered for **high-probability** accesses
  - ▶ **nothing is lost: my daily routine is public**
- I really want to protect **unusual accesses** to documents that might reveal something
  - ▶ I am checking my medical records from some time ago...

$\mathcal{M}$

$\mathcal{O}$

$N$  blocks of data



$B_1$   $B_2$   $B_3$   $\dots$   $B_N$

$N$  encrypted blocks

probability  $p$

$\mathcal{M}$   
-----  
 $\mathcal{O}$

$B_1$   $B_2$   $B_3$   $\dots$   $B_N$

$N$  encrypted blocks

probability  $p$

$\mathcal{M}$

$\mathcal{O}$

Stash

$B_7$   $B_9$   $\dots$   $B_{N-2}$

$pN$  expected blocks



$B_1$   $B_2$   $B_3$   $\dots$   $B_N$

$N$  encrypted blocks

probability  $p$

$\mathcal{M}$

$\mathcal{O}$

Stash

$B_7$   $B_9$   $\dots$   $B_{N-2}$

$pN$  expected blocks

$B_1$   $B_2$   $B_3$   $\dots$   $B_N$

$N$  encrypted blocks

probability  $p$

$\mathcal{M}$

$\mathcal{O}$

Stash

$B_7$   $B_9$   $\dots$   $B_{N-2}$

$pN$  expected blocks

$B_1$   $B_2$   $B_3$   $\dots$   $B_N$

$N$  encrypted blocks

probability  $p$

$\mathcal{M}$

$\mathcal{O}$

Stash

$B_7$   $B_9$   $\dots$   $B_{N-2}$

$pN$  expected blocks

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:

## Overwrite Phase

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it

## Overwrite Phase

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash

## Overwrite Phase

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it

## Overwrite Phase

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:

## Overwrite Phase



# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head
  - ▶ add  $B_i$  to stash

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head
  - ▶ add  $B_i$  to stash
  - ▶ ask  $\mathcal{M}$  for a randomly selected block, re-encrypt it and upload to the same location

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head
  - ▶ add  $B_i$  to stash
  - ▶ ask  $\mathcal{M}$  for a randomly selected block, re-encrypt it and upload to the same location
- if tail

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head
  - ▶ add  $B_i$  to stash
  - ▶ ask  $\mathcal{M}$  for a randomly selected block, re-encrypt it and upload to the same location
- if tail
  - ▶ download  $B_i$  from  $\mathcal{M}$

# Querying for $B_i$

## Download Phase

- if  $B_i$  is found in the stash:
  - ▶ return it
  - ▶ remove it from the stash
  - ▶ ask  $\mathcal{M}$  for a random block and then discard it
- if  $B_i$  is not found in the stash:
  - ▶ ask  $\mathcal{M}$  for  $B_i$

## Overwrite Phase

- Toss a coin with probability  $(p, 1 - p)$
- if head
  - ▶ add  $B_i$  to stash
  - ▶ ask  $\mathcal{M}$  for a randomly selected block, re-encrypt it and upload to the same location
- if tail
  - ▶ download  $B_i$  from  $\mathcal{M}$
  - ▶ decrypt and re-encrypt and upload it to the same location



# Analysis

- Two blocks of communication

# Analysis

- Two blocks of communication
- $p = \omega(\log n/n)$  with  $\omega(\log n)$  client memory

# Analysis

- Two blocks of communication
- $p = \omega(\log n/n)$  with  $\omega(\log n)$  client memory
- $\epsilon = \Theta(\log n)$

# Analysis

- Two blocks of communication
- $p = \omega(\log n/n)$  with  $\omega(\log n)$  client memory
- $\epsilon = \Theta(\log n)$

# Analysis

- Two blocks of communication
- $p = \omega(\log n/n)$  with  $\omega(\log n)$  client memory
- $\epsilon = \Theta(\log n)$

## Theorem

*For any  $\epsilon \geq 0$ , any DP-RAM with error probability  $\alpha \geq 0$  in the balls and bins model and a client that stores at most  $c$  blocks must operate on*

$$\Omega \left( \log_c \left( \frac{(1 - \alpha) \cdot n}{e^\epsilon} \right) \right)$$

*records.*

# Conclusions

- **Leaking the access pattern can be dangerous**

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data



# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications
- Other security notions?

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications
- Other security notions?
- Better analysis under reasonable assumptions for access distributions?  
Zipf?

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications
- Other security notions?
- Better analysis under reasonable assumptions for access distributions?  
Zipf?

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications
- Other security notions?
- Better analysis under reasonable assumptions for access distributions?  
Zipf?

# Conclusions

- **Leaking the access pattern can be dangerous**
- **Obliviousness**: a strong security notion
  - ▶ requires  $> 20\times$  overhead for 4 Giga of data
- **Differential Privacy**: a weaker security notion
  - ▶ very efficient
  - ▶ suitable only for specific applications
- Other security notions?
- Better analysis under reasonable assumptions for access distributions?  
Zipf?

<https://github.com/giuper/talks/nonTechnical/oramTutorial.pdf>



# The original papers



O. Goldreich.

Towards a Theory of Software Protection and Simulation by Oblivious RAMs.

In *STOC*, 1987.



R. Ostrovsky.

Efficient computation on oblivious RAMs.

In *STOC*, pages 514–523, 1990.



O. Goldreich and R. Ostrovsky.

Software Protection and Simulation on Oblivious RAMs.

*J. ACM*, 43(3), 1996.

# Asymptotics



G. Asharov, I. Komargodski, W.-K. Lin, K. Nayak, E. Peserico, and E. Shi.

OptORAMa: Optimal oblivious RAM.

Cryptology ePrint Archive, Report 2018/892.



S. Patel, G. Persiano, M. Raykova, and K. Yeo.

PanORAMa: Oblivious RAM with logarithmic overhead.

Cryptology ePrint Archive, Report 2018/373, 2018.

<https://eprint.iacr.org/2018/373>.

FOCS '18.

## Constant client memory

# Efficient constructions for large blocks



E. Stefanov, M. van Dijk, E. Shi, C. Fletcher, L. Ren, X. Yu, and S. Devadas.

Path ORAM: An Extremely Simple Oblivious RAM Protocol.  
In *CCS '13*, pages 299–310, 2013.



S. Devadas, M. van Dijk, C. W. Fletcher, L. Ren, E. Shi, and D. Wichs.

Onion ORAM: A constant bandwidth blowup oblivious RAM.  
In *TCC*, pages 145–174, 2015.



T. Moataz, T. Mayberry, and E.-O. Blass.

Constant communication ORAM with small blocksize.  
In *CCS*, pages 862–873, 2015.

**polylog(n)-bit blocks**

# Efficient constructions for large client memory



L. Ren, C. Fletcher, A. Kwon, E. Stefanov, E. Shi, M. van Dijk, and S. Devadas.

Constants count: Practical improvements to oblivious RAM.  
In *USENIX Security*, pages 415–430, 2015.



E. Stefanov, E. Shi, and D. X. Song.

Towards practical oblivious RAM.  
In *NDSS 2012*, 2012.





S. Patel, G. Persiano, and K. Yeo.


Recursive ORAMs with practical constructions.  
Cryptology ePrint Archive, Report 2017/964, 2017.


$O(\sqrt{n})$  **client memory**

# Lower bounds

 E. Boyle and M. Naor.  
Is There an Oblivious RAM Lower Bound?  
In *ITCS*, pages 357–368, 2016.

 K. G. Larsen and J. B. Nielsen.  
Yes, there is an oblivious RAM lower bound!  
Cryptology ePrint Archive, Report 2018/423, 2018.  
CRYPTO '18

 M. Weiss and D. Wichs.  
Is there an Oblivious RAM lower bound for online reads?  
Cryptology ePrint Archive, Report 2018/619, 2018.  
TCC '18

 G. Persiano and K. Yeo.  
Lower bounds for differentially private RAMs.  
Cryptology ePrint Archive, Report 2018/1051, 2018.  
Eurocrypt '19

# Oblivious Sorting and Shuffling



M. T. Goodrich.

Zig-Zag Sort: A Simple Deterministic Data-oblivious Sorting Algorithm Running in  $O(N \log N)$  Time.

In *STOC*, pages 684–693, 2014.



O. Ohrimenko, M. T. Goodrich, R. Tamassia, and E. Upfal.

The melbourne shuffle: Improving oblivious storage in the cloud.

In *ICALP '14*, pages 556–567. Springer, 2014.



S. Patel, G. Persiano, and K. Yeo.

CacheShuffle: A Family of Oblivious Shuffles.

In *ICALP*, pages 161:1–161:13, 2018.